



REGRESSION

▼ IMPORT

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import mean_absolute_error, mean_squared_error
from joblib import dump, load
import scipy as sp
df = pd.read_csv("Advertising.csv")
```

```
1 ✓ import numpy as np
2   import pandas as pd
3   import matplotlib.pyplot as plt
4   import seaborn as sns
5
6   from sklearn.model_selection import train_test_split
7   from sklearn.linear_model import LinearRegression
8   from sklearn.preprocessing import PolynomialFeatures
9   from sklearn.metrics import mean_absolute_error, mean_squared_error
10  from joblib import dump, load
11  import scipy as sp
✓ 2.2s

1 df = pd.read_csv("Advertising.csv")
✓ 0.6s
```

▼ Simple Linear Regr


- `sns.regplot(data=df, x="total_spend", y="sales")` : regresyon plotu çizer

```
1 plt.figure(figsize=(12,7), dpi=120)
2 sns.regplot(data=df, x="total_spend", y="sales")
```

✓ 0.5s

Python

<AxesSubplot:xlabel='total_spend', ylabel='sales'>



- Seaborn kodu ile elle hesaplanan poly fit komutunun karşılaştırılması

```
1 x = df["total_spend"]
2 y = df["sales"]
```

✓ 0.1s

```
1 np.polyfit(x,y,deg=1)
```


✓ 0.5s

```
array([0.04868788, 4.24302822])
```

```
1 potential_spend=np.linspace(0,500,100)
2 predicted_sales = 0.04868788*potential_spend + 4.24302822
3
4 sns.scatterplot(x="total_spend", y="sales", data=df)
5 plt.plot(potential_spend,predicted_sales, color="red")
```

✓ 0.4s

```
[<matplotlib.lines.Line2D at 0x1770cfabf10>]
```



▼ Scikit-Learn

- Import

```
6 from sklearn.model_selection import train_test_split
7 from sklearn.linear_model import LinearRegression
```

- veri setini düzenleme. sonuçların gideceği yeni bir y parametresi oluşturuldu

```
1 x= df.drop("sales", axis=1)
2 y= df["sales"]
3 print(x)
4 print(y)
```

✓ 0.4s

Output exceeds the [size limit](#). Open the full output in the [editor](#)

	TV	radio	newspaper
0	230.1	37.8	69.2
1	44.5	39.3	45.1
2	17.2	45.9	69.3
3	151.5	41.3	58.5
4	180.8	10.8	58.4
..
195	38.2	3.7	13.8
196	94.2	4.9	8.1
197	177.0	9.3	6.4
198	283.6	42.0	66.2
199	232.1	8.6	8.7

[200 rows x 3 columns]

0 22.1

...

197 12.8

198 25.5

199 13.4

Name: sales, Length: 200, dtype: float64

- test ve train model sayıları
test_size=0.3 : veriin %30u test için ayrıldı

random_state=101 : rastgele başlangıç değeri

```
1 x_train, x_test, y_train, y_test = train_test_split(x,
2 y, random_state=101, test_size=0.3)
✓ 0.5s

1 print(len(df))
2 print(len(x_train))
3 print(len(x_test))
✓ 0.1s

200
140
60
```

- model= LinearRegression() : model oluşturma
model.fit(x_train, y_train) : modeli yerleştirme
model.predict(x_test) : öngörülen x değerleri

```
1 model= LinearRegression()
✓ 0.5s Python

1 model.fit(x_train, y_train)
✓ 0.8s Python

LinearRegression()

1 model.predict(x_test)
✓ 0.8s Python

array([15.74131332, 19.61062568, 11.44888935, 17.00819787,  9.17285676,
       7.01248287, 20.28992463, 17.29953992,  9.77584467, 19.22194224,
       12.40503154, 13.89234998, 13.72541098, 21.28794031, 18.42456638,
       9.98198406, 15.55228966,  7.68913693,  7.55614992, 20.40311209,
       7.79215204, 18.24214098, 24.68631904, 22.82199068,  7.97962085,
       12.65207264, 21.46925937,  8.05228573, 12.42315981, 12.50719678,
       10.77757812, 19.24460093, 10.070269  ,  6.70779999, 17.31492147,
       7.76764327,  9.25393336,  8.27834697, 10.58105585, 10.63591128,
       13.01002595,  9.77192057, 10.21469861,  8.04572042, 11.5671075 ,
       10.08368001,  8.99806574, 16.25388914, 13.23942315, 20.81493419,
       12.49727439, 13.96615898, 17.56285075, 11.14537013, 12.56261468,
       5.50870279, 23.29465134, 12.62409688, 18.77399978,
       15.18785675])
```

- `print(df["sales"].mean())` : ortalama mean değeri ile MSE ve RMSE kıyaslanmalı ki hata paylarının ne kadar yüksek olduğu ile alakalı bir sonuç elde edilebilsin. %10 civarı çıktı. fena değil.
MSE toplam noktaların ne kadar uzakta olduğunu
RMSE dataların içindeki alakasız, ayırık noktaların ne kadar çok olduğunu gösterir


```
1 test_predictions = model.predict(x_test)
```

✓ 0.3s

```
1 print(df["sales"].mean())  
2 sns.histplot(data=df, x="sales", bins=20);
```

✓ 0.3s

14.022500000000003



```
1 mean_absolute_error(y_test, test_predictions)  
2 # Mean value ile kıyaslanmalı
```

✓ 0.4s

1.2137457736144808

```
1 np.sqrt(mean_squared_error(y_test, test_predictions))  
2 #root mean square error
```

✓ 0.4s

1.5161519375993877

- Residuals and residual plot


```
1 test_residuals = y_test - test_predictions
2 test_residuals.head()
```


✓ 0.1s

```
37    -1.041313
109     0.189374
31      0.451111
89     -0.308198
66      0.327143
Name: sales, dtype: float64
```

```
1 sns.scatterplot(x=y_test, y=test_residuals)
2 plt.axhline(y= 0, color="red", ls="--")
```

✓ 0.2s

<matplotlib.lines.Line2D at 0x1909c787220>




- residual plot

```

1 # Create a figure and axis to plot on
2 fig, ax = plt.subplots(figsize=(6,4),dpi=100)
3 # probplot returns the raw values if needed
4 # we just want to see the plot, so we assign these values
5 _ = sp.stats.probplot(test_residuals,plot=ax)

```

✓ 0.3s



- Coef
 - 1 birim TV artışı 0.045 birim satış artışı
 - 1 birim RADYO artışı 0.188birim satış artışı
 - 1 birim GAZETE artışı 0 birim satış artışı
 - Gazetenin satışa etkisi yok

```
1 final_model = LinearRegression()
2 final_model.fit(x,y)
3 print(final_model.coef_)
✓ 0.3s

[ 0.04576465  0.18853002 -0.00103749]
```

```
1 x.head()
✓ 0.4s
```

	TV	radio	newspaper
0	230.1	37.8	69.2
1	44.5	39.3	45.1
2	17.2	45.9	69.3
3	151.5	41.3	58.5
4	180.8	10.8	58.4

- `dump(final_model, 'sales_model.joblib')` : sonuçları kaydeder.
`loaded_model = load('sales_model.joblib')` : sonuçları çeker ve gösterilmesine imkan verir

```
1 dump(final_model, 'sales_model.joblib')
✓ 0.1s

['sales_model.joblib']

1 loaded_model = load('sales_model.joblib')
✓ 0.9s

1 loaded_model.coef_
✓ 0.8s

array([ 0.04576465,  0.18853002, -0.00103749])
```

- campaign = [[149,22,12]] : 149 k TV, 22 k radyo, 12 k gazete reklamı verildiğinde loaded_model.predict(campaign) : sonuçta kaç birimlik satış elde edilir

```
1 campaign = [[149,22,12]]
✓ 0.8s

1 loaded_model.predict(campaign)
✓ 0.1s

array([13.893032])
```

▼ Polynomial Regression

- Dataset düzenleme

```
1 x = df.drop("sales", axis=1)
2 y = df["sales"]
3 print(x.head())
4 print(y.head())
```

✓ 0.2s

	TV	radio	newspaper
0	230.1	37.8	69.2
1	44.5	39.3	45.1
2	17.2	45.9	69.3
3	151.5	41.3	58.5
4	180.8	10.8	58.4

	TV
0	22.1
1	10.4
2	9.3
3	18.5
4	12.9

Name: sales, dtype: float64

- polynom özellikleri ekleme

```
1 polynomial_converter = PolynomialFeatures(degree= 2, include_bias= False)
✓ 0.4s

1 polynomial_converter.fit(x)
✓ 0.9s

PolynomialFeatures(include_bias=False)

1 poly_features = polynomial_converter.fit_transform(x)
✓ 0.1s

1 print(polynomial_converter.transform(x).shape)
2 print(x.shape)
✓ 0.9s

(200, 9)
(200, 3)
```

- poly değerler. 1 2 3 dataframeden gelen. 4 5 6 , ilk 3'ün karesi. 7 8 9 birbirleriyle çarpımı

```
1 print(x.iloc[0])
2 print(poly_features[0])
✓ 0.4s

TV          230.1
radio       37.8
newspaper   69.2
Name: 0, dtype: float64
[2.301000e+02 3.780000e+01 6.920000e+01 5.294601e+04 8.697780e+03
 1.592292e+04 1.428840e+03 2.615760e+03 4.788640e+03]
```

- polynomial sonuçlar, 2. dereceden

```
1 x_train, x_test, y_train, y_test = train_test_split(poly_features,  
2 | y, test_size=0.3, random_state=101)
```

✓ 0.1s

```
1 model = LinearRegression()  
2 model.fit(x_train, y_train)
```

✓ 0.2s

LinearRegression()

```
1 test_predictions = model.predict(x_test)
```

✓ 0.2s

```
1 MAE = mean_absolute_error(y_test, test_predictions)  
2 MSE = mean_squared_error(y_test, test_predictions)  
3 RMSE = np.sqrt(MSE)
```

✓ 0.5s

```
1 print(f"MAE = {round(MAE,3)}")  
2 print(f"MSE = {round(MSE,3)}")  
3 print(f"RMSE = {round(RMSE,3)}")
```

✓ 0.1s

MAE = 0.49
MSE = 0.442
RMSE = 0.665

▼ Model Seçimi

- kaçıncı dereceden model seçileceğini belirlemek için min. error rate veren bir derece seçilmeli

```

1 # TRAINING ERROR PER DEGREE
2 train_rmse_errors = []
3 # TEST ERROR PER DEGREE
4 test_rmse_errors = []
5
6 for d in range(1,10):
7
8     # CREATE POLY DATA SET FOR DEGREE "d"
9     polynomial_converter = PolynomialFeatures(degree=d,include_bias=False)
10    poly_features = polynomial_converter.fit_transform(x)
11
12    # SPLIT THIS NEW POLY DATA SET
13    x_train, x_test, y_train, y_test = train_test_split(poly_features,
14    | y, test_size=0.3, random_state=101)
15
16    # TRAIN ON THIS NEW POLY SET
17    model = LinearRegression(fit_intercept=True)
18    model.fit(x_train,y_train)
19
20    # PREDICT ON BOTH TRAIN AND TEST
21    train_pred = model.predict(x_train)
22    test_pred = model.predict(x_test)
23
24    # Calculate Errors
25
26    # Errors on Train Set
27    train_RMSE = np.sqrt(mean_squared_error(y_train,train_pred))
28
29    # Errors on Test Set
30    test_RMSE = np.sqrt(mean_squared_error(y_test,test_pred))
31
32    # Append errors to lists for plotting later
33
34    train_rmse_errors.append(train_RMSE)
35    test_rmse_errors.append(test_RMSE)

```

✓ 0.1s

Python

- 1-9 dereceden polinomların sonuçları


```
1 train_rmse_errors
✓ 0.3s

[1.734594124329376,
 0.5879574085292233,
 0.4339344356902067,
 0.35170836883993534,
 0.2509342952029336,
 0.19933332834273104,
 5.4214215994181805,
 0.14237972100695595,
 0.16675080548552418]

1 test_rmse_errors
✓ 0.3s


[1.5161519375993873,
 0.6646431757269196,
 0.5803286825231453,
 0.5077742624232109,
 2.5758247603435955,
 4.490868529265006,
 1381.404235838588,
 4449.5681972303655,
 95893.0265813161]
```

- NOT : Kırılmanın olduğu noktanın bir öncesinden alırsan daha iyi olur. 4. derecede nerede overfit yaptığını bilemeyebilirsin.
(Bu örnek için)

```
1 plt.plot(range(1,6),train_rmse_errors[:5],label='TRAIN')
2 plt.plot(range(1,6),test_rmse_errors[:5],label='TEST')
3 plt.xlabel("Polynomial Complexity")
4 plt.ylabel("RMSE")
5 plt.legend()
```

✓ 0.3s

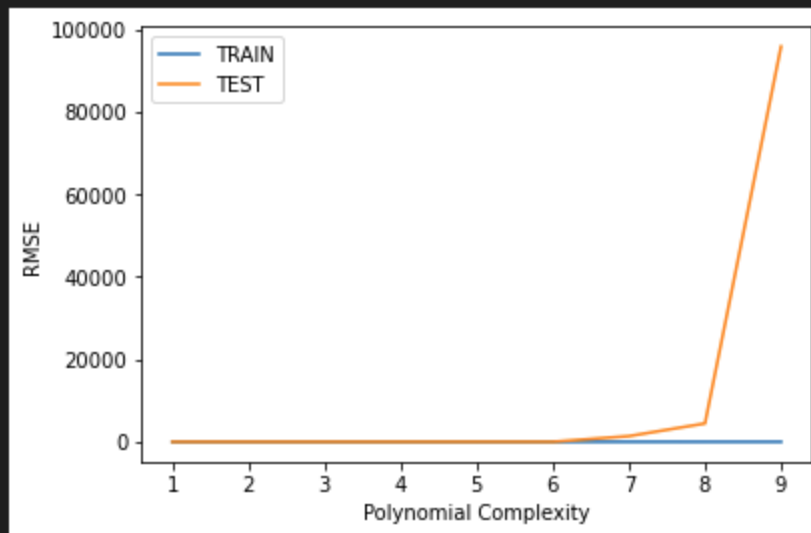
<matplotlib.legend.Legend at 0x208a461d730>



```
1 plt.plot(range(1,10),train_rmse_errors,label='TRAIN')
2 plt.plot(range(1,10),test_rmse_errors,label='TEST')
3 plt.xlabel("Polynomial Complexity")
4 plt.ylabel("RMSE")
5 plt.legend()
```

✓ 0.3s

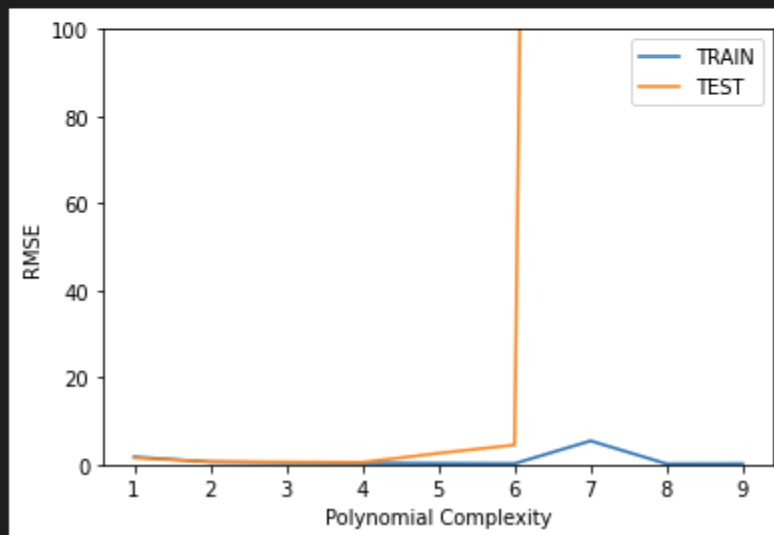
<matplotlib.legend.Legend at 0x208a46a4940>



```
1 plt.plot(range(1,10),train_rmse_errors,label='TRAIN')
2 plt.plot(range(1,10),test_rmse_errors,label='TEST')
3 plt.xlabel("Polynomial Complexity")
4 plt.ylabel("RMSE")
5 plt.ylim(0,100)
6 plt.legend()
```

✓ 0.1s

<matplotlib.legend.Legend at 0x208a466cc70>



▼ Regularizations

- regularization

```

1 df = pd.read_csv("Advertising.csv")
2 X = df.drop("sales", axis=1)
3 y = df["sales"]
4 print(X)
5 print(y)

```

✓ 0.6s

Output exceeds the [size limit](#). Open the full output

	TV	radio	newspaper
0	230.1	37.8	69.2
1	44.5	39.3	45.1
2	17.2	45.9	69.3
3	151.5	41.3	58.5
4	180.8	10.8	58.4
...
195	38.2	3.7	13.8
196	94.2	4.9	8.1
197	177.0	9.3	6.4
198	283.6	42.0	66.2
199	232.1	8.6	8.7

[200 rows x 3 columns]

0 22.1

...

197 12.8

198 25.5

199 13.4

Name: sales, Length: 200, dtype: float64

```
1 polynomial_converter = PolynomialFeatures(degree=3, include_bias=False)
✓ 0.7s
```

```
1 poly_features = polynomial_converter.fit_transform(X)
✓ 0.5s
```

```
1 X_train, X_test, y_train, y_test = train_test_split(poly_features,
2 | y, test_size=0.3, random_state=101)
✓ 0.7s
```

```
1 scaler = StandardScaler()
2 scaler.fit(X_train)
✓ 0.1s
```

StandardScaler()

```
1 X_train = scaler.transform(X_train)
2 X_test = scaler.transform(X_test)
✓ 0.3s
```

- Ridge regression

```
1 from sklearn.linear_model import Ridge
✓ 0.6s

1 ridge_model = Ridge(alpha=10)
2 ridge_model.fit(X_train,y_train)
✓ 0.3s

Ridge(alpha=10)

1 test_predictions = ridge_model.predict(X_test)
✓ 0.1s

1 MAE = mean_absolute_error(y_test,test_predictions)
2 MSE = mean_squared_error(y_test,test_predictions)
3 RMSE = np.sqrt(MSE)
4 print(f"MAE = {round(MAE,3)}")
5 print(f"MSE = {round(MSE,3)}")
6 print(f"RMSE = {round(RMSE,3)}")
✓ 0.7s

MAE = 0.577
MSE = 0.8
RMSE = 0.895
```

- Ridge Cross Validation

```
1 from sklearn.linear_model import RidgeCV
✓ 0.5s Pyt

1 ridge_cv_model = RidgeCV(alphas = (0.1, 1, 10), scoring="neg_mean_absolute_error")
✓ 0.1s Pyt

1 ridge_cv_model.fit(X_train, y_train)
✓ 0.7s Pyt
RidgeCV(alphas=array([ 0.1, 1. , 10. ]), scoring='neg_mean_absolute_error')

1 from sklearn.metrics import SCORERS
✓ 0.5s Pyt

1 test_predictions = ridge_cv_model.predict(X_test)
✓ 0.4s Pyt

1 MAE = mean_absolute_error(y_test, test_predictions)
2 MSE = mean_squared_error(y_test, test_predictions)
3 RMSE = np.sqrt(MSE)
4 print(f"MAE = {round(MAE, 3)}")
5 print(f"MSE = {round(MSE, 3)}")
6 print(f"RMSE = {round(RMSE, 3)}")
✓ 0.1s Pyt

MAE = 0.427
MSE = 0.382
RMSE = 0.618
```

- Lasso Regression


```
1 from sklearn.linear_model import LassoCV
✓ 0.7s

1 Lasso_cv_model = LassoCV(eps=0.1, n_alphas=100, cv=5)
✓ 0.9s

1 Lasso_cv_model.fit(X_train, y_train)
✓ 0.3s

LassoCV(cv=5, eps=0.1)

1 test_predictions = Lasso_cv_model.predict(X_test)
✓ 0.9s

1 MAE = mean_absolute_error(y_test, test_predictions)
2 MSE = mean_squared_error(y_test, test_predictions)
3 RMSE = np.sqrt(MSE)
4 print(f"MAE = {round(MAE, 3)}")
5 print(f"MSE = {round(MSE, 3)}")
6 print(f"RMSE = {round(RMSE, 3)}")
✓ 0.7s

MAE = 0.654
MSE = 1.279
RMSE = 1.131
```

- Elastic Net

```
1 from sklearn.linear_model import ElasticNetCV
✓ 0.5s

1 elastic_model = ElasticNetCV(l1_ratio=[.1, .5, .7, .9, .95, .99, 1],tol=0.01)
✓ 0.8s

1 elastic_model.fit(X_train,y_train)
✓ 0.6s
ElasticNetCV(l1_ratio=[0.1, 0.5, 0.7, 0.9, 0.95, 0.99, 1], tol=0.01)

1 test_predictions = elastic_model.predict(X_test)
✓ 0.2s

1 MAE = mean_absolute_error(y_test,test_predictions)
2 MSE = mean_squared_error(y_test,test_predictions)
3 RMSE = np.sqrt(MSE)
4 print(f"MAE = {round(MAE,3)}")
5 print(f"MSE = {round(MSE,3)}")
6 print(f"RMSE = {round(RMSE,3)}")
✓ 0.2s
MAE = 0.566
MSE = 0.56
RMSE = 0.749

1 train_predictions = elastic_model.predict(X_train)
2 MAE = mean_absolute_error(y_train,train_predictions)
3 MAE
✓ 0.1s
0.4307582990472369
```

-

▼ Appendix

- Transform

```

1 neighborhood["Lot Frontage"].transform(lambda value: value.fillna(value.mean()))
2

```

✓ 0.7s

Python

- `.apply(str)` : Stringe çevirme

```

1 df["MS SubClass"] = df["MS SubClass"].apply(str)

```

✓ 0.2s

- `pd.get_dummies(direction)` : string veriyi dummy değişkene çevirir

```

1 direction = pd.Series(["up", "up", "down", "down"])
2 direction

```

✓ 0.1s

```

0      up
1      up
2     down
3     down
dtype: object

```

```

1 pd.get_dummies(direction)

```

✓ 0.5s

	down	up
0	0	1
1	0	1
2	1	0
3	1	0

- `pd.get_dummies(direction, drop_first=True)` : ilk kolonu atar

```
1 pd.get_dummies(direction, drop_first=True)
```

✓ 0.7s

	up
0	1
1	1
2	0
3	0

- `df.select_dtypes(include="object")` : Türü obje olan verileri getirir

```
1 df.select_dtypes(include="object")
```

✓ 0.6s

	MS SubClass	MS Zoning	Street	Lot Shape	Land Contour	Utilities	Lot Config	Land Slope
0	20	RL	Pave	IR1	Lvl	AllPub	Corner	Gtl
1	20	RH	Pave	Reg	Lvl	AllPub	Inside	Gtl
2	20	RL	Pave	IR1	Lvl	AllPub	Corner	Gtl
3	20	RL	Pave	Reg	Lvl	AllPub	Corner	Gtl
4	60	RL	Pave	IR1	Lvl	AllPub	Inside	Gtl
...
2925	80	RL	Pave	IR1	Lvl	AllPub	CulDSac	Gtl
2926	20	RL	Pave	IR1	Low	AllPub	Inside	Mod
2927	85	RL	Pave	Reg	Lvl	AllPub	Inside	Gtl
2928	20	RL	Pave	Reg	Lvl	AllPub	Inside	Mod
2929	60	RL	Pave	Reg	Lvl	AllPub	Inside	Mod

- dummy variable oluşturur

```
1 my_object_df = df.select_dtypes(include="object")
✓ 0.1s

1 my_numeric_df = df.select_dtypes(exclude="object")
✓ 0.8s

1 df_object_dummies = pd.get_dummies(my_object_df, drop_first=True)
2 df_object_dummies
✓ 0.2s
```

	MS SubClass_150	MS SubClass_160	MS SubClass_180	MS SubClass_190	MS SubClass_20	Sul
0	0	0	0	0	1	
1	0	0	0	0	1	
2	0	0	0	0	1	
3	0	0	0	0	1	
4	0	0	0	0	0	

- `pd.concat` : data frameleri birleştirir

```

1 final_df = pd.concat([my_numeric_df, df_object_dummies], axis=1)
2 final_df

```

✓ 0.1s

	Lot Frontage	Lot Area	Overall Qual	Overall Cond	Year Built	Year Remod/Add	Mas Vnr Area	BsmtFin SF 1
0	141.000000	31770	6	5	1960	1960	112.0	639.0
1	80.000000	11622	5	6	1961	1961	0.0	468.0
2	81.000000	14267	6	6	1958	1958	108.0	923.0
3	93.000000	11160	7	5	1968	1968	0.0	1065.0
4	74.000000	13830	5	5	1997	1998	0.0	791.0
...
2925	37.000000	7937	6	6	1984	1984	0.0	819.0
2926	75.144444	8885	5	5	1983	1983	0.0	301.0
2927	62.000000	10441	5	5	1992	1992	0.0	337.0
2928	77.000000	10010	5	5	1974	1975	0.0	1071.0
2929	74.000000	9627	7	5	1993	1994	94.0	758.0

2925 rows × 274 columns

▼ Grid Search

- Import and prepare

```

1 X_train,X_test, y_train,y_test = train_test_split(X,y,test_size=0.3, random_state=101)
2 scaler = StandardScaler()
3 scaler.fit(X_train)
4 X_train = scaler.transform(X_train)
5 X_test = scaler.transform(X_test)

```

✓ 0.1s

```

1 from sklearn.linear_model import ElasticNet

```

✓ 0.4s

- grid search cv

```
1 base_elastic_net_model = ElasticNet()
✓ 0.8s

1 param_grid = {
2     "alpha": [0.1, 1, 5, 10, 50, 100],
3     "l1_ratio": [.1, .5, .7, .95, .99, 1]
4 }
✓ 0.5s

1 from sklearn.model_selection import GridSearchCV
✓ 0.8s

1 grid_model = GridSearchCV(
2     estimator = base_elastic_net_model,
3     param_grid = param_grid,
4     scoring = "neg_mean_squared_error",
5     cv = 5, verbose = 2
6 )
✓ 0.8s

1 grid_model.fit(X_train, y_train)
✓ 0.6s
```

- fitted model.

```
1 grid_model.fit(X_train,y_train)
✓ 0.6s

Output exceeds the size limit. Open the full output data in a text editor
Fitting 5 folds for each of 36 candidates, totalling 180 fits
[CV] END .....alpha=0.1, l1_ratio=0.1; total time= 0.0s
[CV] END .....alpha=0.1, l1_ratio=0.1; total time= 0.0s
[CV] END .....alpha=0.1, l1_ratio=0.1; total time= 0.0s
[CV] END .....alpha=0.1, l1_ratio=0.1; total time= 0.0s
[CV] END .....alpha=0.1, l1_ratio=0.1; total time= 0.0s
[CV] END .....alpha=0.1, l1_ratio=0.5; total time= 0.0s
[CV] END .....alpha=0.1, l1_ratio=0.5; total time= 0.0s
[CV] END .....alpha=0.1, l1_ratio=0.5; total time= 0.0s
[CV] END .....alpha=0.1, l1_ratio=0.5; total time= 0.0s
[CV] END .....alpha=0.1, l1_ratio=0.5; total time= 0.0s
[CV] END .....alpha=0.1, l1_ratio=0.7; total time= 0.0s
[CV] END .....alpha=0.1, l1_ratio=0.7; total time= 0.0s
[CV] END .....alpha=0.1, l1_ratio=0.7; total time= 0.0s
[CV] END .....alpha=0.1, l1_ratio=0.7; total time= 0.0s
...
[CV] END .....alpha=100, l1_ratio=1; total time= 0.0s
[CV] END .....alpha=100, l1_ratio=1; total time= 0.0s
[CV] END .....alpha=100, l1_ratio=1; total time= 0.0s
[CV] END .....alpha=100, l1_ratio=1; total time= 0.0s

GridSearchCV(cv=5, estimator=ElasticNet(),
             param_grid={'alpha': [0.1, 1, 5, 10, 50, 100],
                         'l1_ratio': [0.1, 0.5, 0.7, 0.95, 0.99, 1]},
             scoring='neg_mean_squared_error', verbose=2)
```

- en iyi modeli seçmek için
farklı alpha ve l1_ratio değerleri için denendi ve en iyisi bulundu.

```
1 grid_model.best_estimator_
✓ 0.5s

ElasticNet(alpha=0.1, l1_ratio=1)
```

- Aynısı ama dictionary formunda


```
1 grid_model.best_params_  
✓ 0.4s  
{'alpha': 0.1, 'l1_ratio': 1}
```

▼ Project Overview

- project data import

```
1 df = pd.read_csv("../00 DATA/AMES_Final_DF.csv")  
✓ 0.1s
```

```
1 df.head()  
✓ 0.9s
```

	Lot Frontage	Lot Area	Overall Qual	Overall Cond	Year Built	Year Remod/Add	Mas Vnr Area
0	141.0	31770	6	5	1960	1960	112.0
1	80.0	11622	5	6	1961	1961	0.0
2	81.0	14267	6	6	1958	1958	108.0
3	93.0	11160	7	5	1968	1968	0.0
4	74.0	13830	5	5	1997	1998	0.0

5 rows × 274 columns

- data and train test set preparing

```
1 X = df.drop("SalePrice", axis=1)
2 y = df["SalePrice"]
✓ 0.1s Python

1 from sklearn.model_selection import train_test_split
✓ 0.8s Python

1 X_train,X_test, y_train,y_test = train_test_split(X,y,test_size=0.1, random_state=101)
2
✓ 0.7s Python

1 from sklearn.preprocessing import StandardScaler
✓ 0.9s Python

1 scaler = StandardScaler()
2 scaler.fit(X_train)
3 scaled_X_train = scaler.transform(X_train)
4
5 # scaled_X_train = scaler.fit_transform(X_train)
6 # Üsteki 2 ve 3. satırın aynısını yapar
7
8 scaled_X_test = scaler.transform(X_test)
✓ 0.1s Python
```

- elasticnet and parameters

```
1 from sklearn.linear_model import ElasticNet
✓ 0.3s

1 base_elastic_model = ElasticNet()
✓ 0.1s

1 param_grid = {
2     "alpha": [0.1, 1, 5, 10, 50, 100],
3     "l1_ratio": [.1, .5, .7, .95, .99, 1]
4 }
✓ 0.4s
```

- Grid Search model and best parameters

```
1 from sklearn.model_selection import GridSearchCV
✓ 0.4s

1 grid_model = GridSearchCV(
2     base_elastic_model,
3     param_grid= param_grid,
4     scoring="neg_mean_squared_error",
5     cv=5, verbose=1
6 )
✓ 0.6s

1 grid_model.fit(scaled_X_train,y_train)
✓ 4m 41.6s

Fitting 5 folds for each of 36 candidates, totalling 180 fits

GridSearchCV(cv=5, estimator=ElasticNet(max_iter=1000000),
             param_grid={'alpha': [0.1, 1, 5, 10, 50, 100],
                         'l1_ratio': [0.1, 0.5, 0.7, 0.95, 0.99, 1]},
             scoring='neg_mean_squared_error', verbose=1)

1 | grid_model.best_params_
✓ 0.8s

{'alpha': 100, 'l1_ratio': 1}
```

- mean squared error & mean absolute error

```
1 y_pred = grid_model.predict(scaled_X_test)
```

✓ 0.4s

```
1 from sklearn.metrics import mean_squared_error, mean_absolute_error
```

✓ 0.4s

```
1 mean_absolute_error(y_test,y_pred)
```

✓ 0.5s

14195.354900562168

```
1 np.sqrt(mean_squared_error(y_test,y_pred))
```

✓ 0.1s

20558.508566893164