



# Random Forest

## ▼ PENGUIN SPECIES

- INTRO

### Imports

+ Code + Markdown

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
```

✓ 3.5s

```
1 df = pd.read_csv("penguins_size.csv")
```

✓ 0.6s

```
1 df = df.dropna()
2 df.head()
```

✓ 0.1s

	species	island	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	sex
0	Adelie	Torgersen	39.1	18.7	181.0	3750.0	MALE
1	Adelie	Torgersen	39.5	17.4	186.0	3800.0	FEMALE
2	Adelie	Torgersen	40.3	18.0	195.0	3250.0	FEMALE
4	Adelie	Torgersen	36.7	19.3	193.0	3450.0	FEMALE
5	Adelie	Torgersen	39.3	20.6	190.0	3650.0	MALE

# Train test split

```
1 X = pd.get_dummies(df.drop("species",axis=1), drop_first=True)
2 y = df["species"]
```

✓ 0.6s

```
1 from sklearn.model_selection import train_test_split
```

✓ 1.4s

```
1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=101)
```

✓ 0.4s

```
1 from sklearn.ensemble import RandomForestClassifier
```

0.3s

```
1 rfc = RandomForestClassifier(n_estimators=10, random_state=101, max_features="auto")
```

0.5s

```
1 rfc.fit(X_train, y_train)
```

0.1s

```
RandomForestClassifier(n_estimators=10, random_state=101)
```

```
1 preds = rfc.predict(X_test)
```

```
2 preds
```

0.1s

```
array(['Chinstrap', 'Gentoo', 'Adelie', 'Chinstrap', 'Gentoo',
       'Chinstrap', 'Adelie', 'Gentoo', 'Chinstrap', 'Gentoo', 'Adelie',
       'Adelie', 'Adelie', 'Gentoo', 'Gentoo', 'Adelie', 'Gentoo',
       'Adelie', 'Adelie', 'Adelie', 'Gentoo', 'Chinstrap', 'Adelie',
       'Adelie', 'Adelie', 'Adelie', 'Chinstrap', 'Gentoo', 'Adelie',
       'Chinstrap', 'Gentoo', 'Chinstrap', 'Gentoo', 'Adelie', 'Adelie',
       'Chinstrap', 'Adelie', 'Gentoo', 'Chinstrap', 'Gentoo', 'Adelie',
       'Adelie', 'Gentoo', 'Adelie', 'Adelie', 'Chinstrap', 'Chinstrap',
       'Chinstrap', 'Chinstrap', 'Chinstrap', 'Adelie', 'Adelie',
```

- Evaluation

# Evaluation

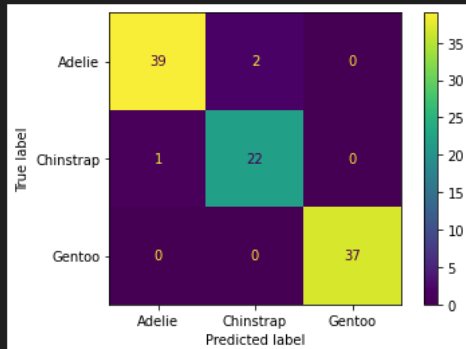
```
1 from sklearn.metrics import confusion_matrix, classification_report, plot_confusion_matrix
```

✓ 0.7s

```
1 plot_confusion_matrix(rfc, X_test, y_test)
```

✓ 0.4s

<sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay at 0x1fb7c723190>



```
1 print(classification_report(y_test, preds))
```

✓ 0.1s

	precision	recall	f1-score	support
Adelie	0.97	0.95	0.96	41
Chinstrap	0.92	0.96	0.94	23
Gentoo	1.00	1.00	1.00	37
accuracy			0.97	101
macro avg	0.96	0.97	0.97	101
weighted avg	0.97	0.97	0.97	101

## ▼ Banknote Authentication

# Imports

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
```

✓ 6.3s

```
1 df = pd.read_csv("data_banknote_authentication.csv")
```

✓ 0.6s

# Data and Viz

```
1 df.head()
```

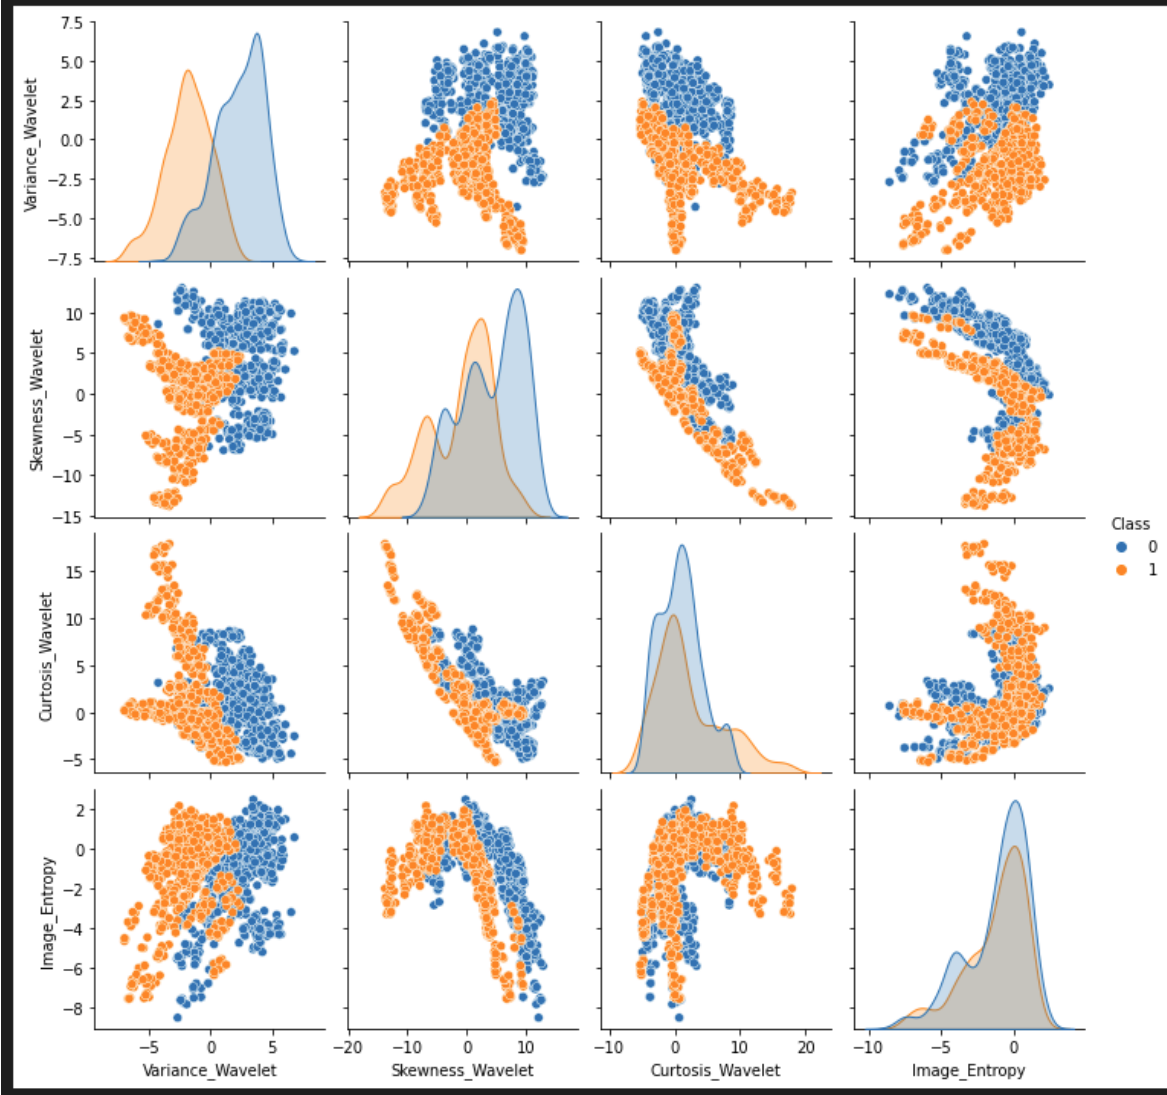
✓ 0.6s

	Variance_Wavelet	Skewness_Wavelet	Curtosis_Wavelet	Image_Entropy	Class
0	3.62160	8.6661	-2.8073	-0.44699	0
1	4.54590	8.1674	-2.4586	-1.46210	0
2	3.86600	-2.6383	1.9242	0.10645	0
3	3.45660	9.5228	-4.0112	-3.59440	0
4	0.32924	-4.4552	4.5718	-0.98880	0

```
1 sns.pairplot(df,hue="Class")
```

✓ 6.5s

<seaborn.axisgrid.PairGrid at 0x19752fd5fa0>



# Train test

```
1 X = df.drop("Class", axis=1)
2 y = df["Class"]
```

✓ 0.4s

```
1 from sklearn.model_selection import train_test_split
```

✓ 0.3s

```
1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.15, random_state=101)
```

✓ 0.4s

```
1 from sklearn.model_selection import GridSearchCV
```

✓ 0.6s

```
1 from sklearn.ensemble import RandomForestClassifier
```

✓ 0.5s

```
1 n_estimators = [64,100,128,200]
2 max_features = [2,3,4]
3 bootstrap = [True, False]
4 oob_score = [True, False]
```

```
1 param_grid = {
2     "n_estimators" : n_estimators,
3     "max_features" : max_features,
4     "bootstrap" : bootstrap,
5     "oob_score" : oob_score
6 }
```

⌵ ⏏

⊗ 0.1s

```
1 rfc = RandomForestClassifier()
```

✓ 0.1s

```
1 grid = GridSearchCV(rfc,param_grid)
```

✓ 0.8s

```
1 grid.fit(X_train, y_train)
```

✓ 1m 15.4s

```
1 grid.best_params_
✓ 0.6s
{'bootstrap': True, 'max_features': 2, 'n_estimators': 64, 'oob_score': False}

1 rfc = RandomForestClassifier(max_features=2, n_estimators=200, oob_score=True)
✓ 0.1s

1 rfc.fit(X_train,y_train)
✓ 1.5s
RandomForestClassifier(max_features=2, n_estimators=200, oob_score=True)

1 rfc.oob_score_
✓ 0.9s
0.9948542024013722
```



# Evaluation

```
1 from sklearn.metrics import classification_report, plot_confusion_matrix
```

✓ 0.5s

```
1 print(classification_report(y_test, predictions))
```


✓ 0.1s

	precision	recall	f1-score	support
0	1.00	0.98	0.99	124
1	0.98	1.00	0.99	82
accuracy			0.99	206
macro avg	0.99	0.99	0.99	206
weighted avg	0.99	0.99	0.99	206

```
1 plot_confusion_matrix(rfc, X_test, y_test)
```

✓ 0.8s

<sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay at 0x197544dcd00



```

1 errors = []
2 misclass = []
3
4 ✓ for n in range(1,200):
5     rfc = RandomForestClassifier(n_estimators=n, max_features=2)
6     rfc.fit(X_train,y_train)
7     preds = rfc.predict(X_test)
8     err = 1 - accuracy_score(y_test,preds)
9     n_missed = np.sum(preds != y_test)
10    # son işlem öngörülen değerlerden kaç tanesinin y_test içindeki değerlerle
11    # ... eşleşmediğini verir. Bu şekilde hatalı öngörü sayısı belirlenebilir.
12    # preds != y_test <<< Predictionlar y_test'e eşit değilse demek
13
14    errors.append(err)
15    misclass.append(n_missed)
16
17


```

✓ 1m 7.6s

```
1 plt.plot(range(1,200),errors)
```

✓ 0.5s


[<matplotlib.lines.Line2D at 0x197556a4880>]



```
1 plt.plot(range(1,200),misclass)
```

✓ 0.5s

[<matplotlib.lines.Line2D at 0x19755705ca0>]



## ▼ Rock Tunnelling

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
```

```
1 df = pd.read_csv("rock_density_xray.csv")
```

# Data and Viz

[+ Code](#)[+ M](#)

```
1 df.head()
```

✓ 0.5s

	Rebound Signal Strength nHz	Rock Density kg/m3
0	72.945124	2.456548
1	14.229877	2.601719
2	36.597334	1.967004
3	9.578899	2.300439
4	21.765897	2.452374


```
1 df.columns=["Signal","Density"]
```

✓ 0.1s

```
1 sns.scatterplot(data = df, x="Signal",y="Density")
```

✓ 0.5s

<AxesSubplot:xlabel='Signal', ylabel='Density'>



# Train Test Split

```
1 X = df["Signal"].values.reshape(-1,1)#değerleri uydurmak için reshape edildi.  
2 # aksi halde hata veriyor  
3 y = df["Density"]
```

✓ 0.4s

Python

```
1 from sklearn.model_selection import train_test_split
```

✓ 0.9s

Python

```
1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, random_state=101)
```

✓ 0.6s

Python

```
1 from sklearn.linear_model import LinearRegression
```

✓ 0.1s

Python

```
1 lr_model = LinearRegression()
```

✓ 0.3s

+ Code

+ Markdown

```
1 lr_model.fit(X_train,y_train)
```

✓ 0.9s

LinearRegression()

```
1 lr_preds = lr_model.predict(X_test)
```

✓ 0.9s

```
1 lr_preds
```

✓ 0.6s

```
array([2.22029657, 2.22047771, 2.22035637, 2.22034337, 2.22039737,  
       2.22050555, 2.22042659, 2.22028877, 2.22034673, 2.22029714,  
       2.22041506, 2.22050153, 2.22043891, 2.22042003, 2.22047022,  
       2.22032403, 2.22033377, 2.22030628, 2.22035154, 2.22035373,  
       2.22029266, 2.22036798, 2.22033018, 2.22030611, 2.22042754,  
       2.22044019, 2.2204142 , 2.22040303, 2.22048946, 2.22047495])
```

```

1 from sklearn.metrics import mean_absolute_error, mean_squared_error
4] ✓ 0.5s

1 mean_absolute_error(y_test,lr_preds)
5] ✓ 0.7s

0.211198973318633

1 np.sqrt(mean_squared_error(y_test,lr_preds))
6] ✓ 0.7s

0.2570051996584629

```

```

1 signal_range = np.arange(0,100)
2 signal_range
] ✓ 0.5s

array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
        17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
        34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50,
        51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67,
        68, 69, 70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81, 82, 83, 84,
        85, 86, 87, 88, 89, 90, 91, 92, 93, 94, 95, 96, 97, 98, 99])

1 signal_pred = lr_model.predict(signal_range.reshape(-1,1))
2 signal_pred
3 # Bu işlem linear regresyonun neden hatalı sonuç verdiğini anlamak için yapıldı.
4 # Grafik sinüs dalgası gibi görünüyor bu sebeple uydurulan doğru
5 # ... direkt bir ortalama değer alıyor ve ortalama üzerinden çalışıyor.
] ✓ 0.2s

Output exceeds the size limit. Open the full output data in a text editor
array([2.22028446, 2.22028673, 2.22028899, 2.22029126, 2.22029353,
        2.22029579, 2.22029806, 2.22030032, 2.22030259, 2.22030485,
        2.22030712, 2.22030938, 2.22031165, 2.22031391, 2.22031618,
        2.22031844, 2.22032071, 2.22032297, 2.22032524, 2.2203275 ,
        2.22032977, 2.22033204, 2.2203343 , 2.22033657, 2.22033883,
        2.2203411 , 2.22034336, 2.22034563, 2.22034789, 2.22035016,
        2.22035242, 2.22035469, 2.22035695, 2.22035922, 2.22036148,
        2.22036375, 2.22036602, 2.22036828, 2.22037055, 2.22037281,
        2.22037508, 2.22037734, 2.22037961, 2.22038187, 2.22038414,
        2.2203864 , 2.22038867, 2.22039093, 2.2203932 , 2.22039546,
        2.22039773, 2.22039999, 2.22040226, 2.22040453, 2.22040679,
        2.22040906, 2.22041132, 2.22041359, 2.22041585, 2.22041812,


```

```
1 plt.figure(figsize=(12,8), dpi=200)
2 sns.scatterplot(data = df, x="Signal",y="Density")
3 plt.plot(signal_range,signal_pred)
```

✓ 0.4s

Python

[<matplotlib.lines.Line2D at 0x216a0907c10>]





```
1 def run_model(model,X_train,y_train,X_test,y_test):
2     #Farklı modellerle test edip sonuçları karşılaştırmak için
3
4     # FIT MODEL
5     model.fit(X_train,y_train)
6
7     # GET METRICS
8     preds = model.predict(X_test)
9     rmse = np.sqrt(mean_squared_error(y_test,preds))
10    mae = mean_absolute_error(y_test,preds)
11    print(f"RMSE : {rmse}")
12    print(f"MAE : {mae}")
13
14    # PLOT THR RESULTS
15    signal_range = np.arange(0,100)
16    signal_pred = model.predict(signal_range.reshape(-1,1))
17
18    plt.figure(figsize=(12,8), dpi=200)
19    sns.scatterplot(data = df, x="Signal",y="Density", color="black")
20
21    plt.plot(signal_range,signal_pred)
```

✓ 0.5s Python


```
1 model = LinearRegression()  
2 run_model(model,X_train,y_train,X_test,y_test)
```

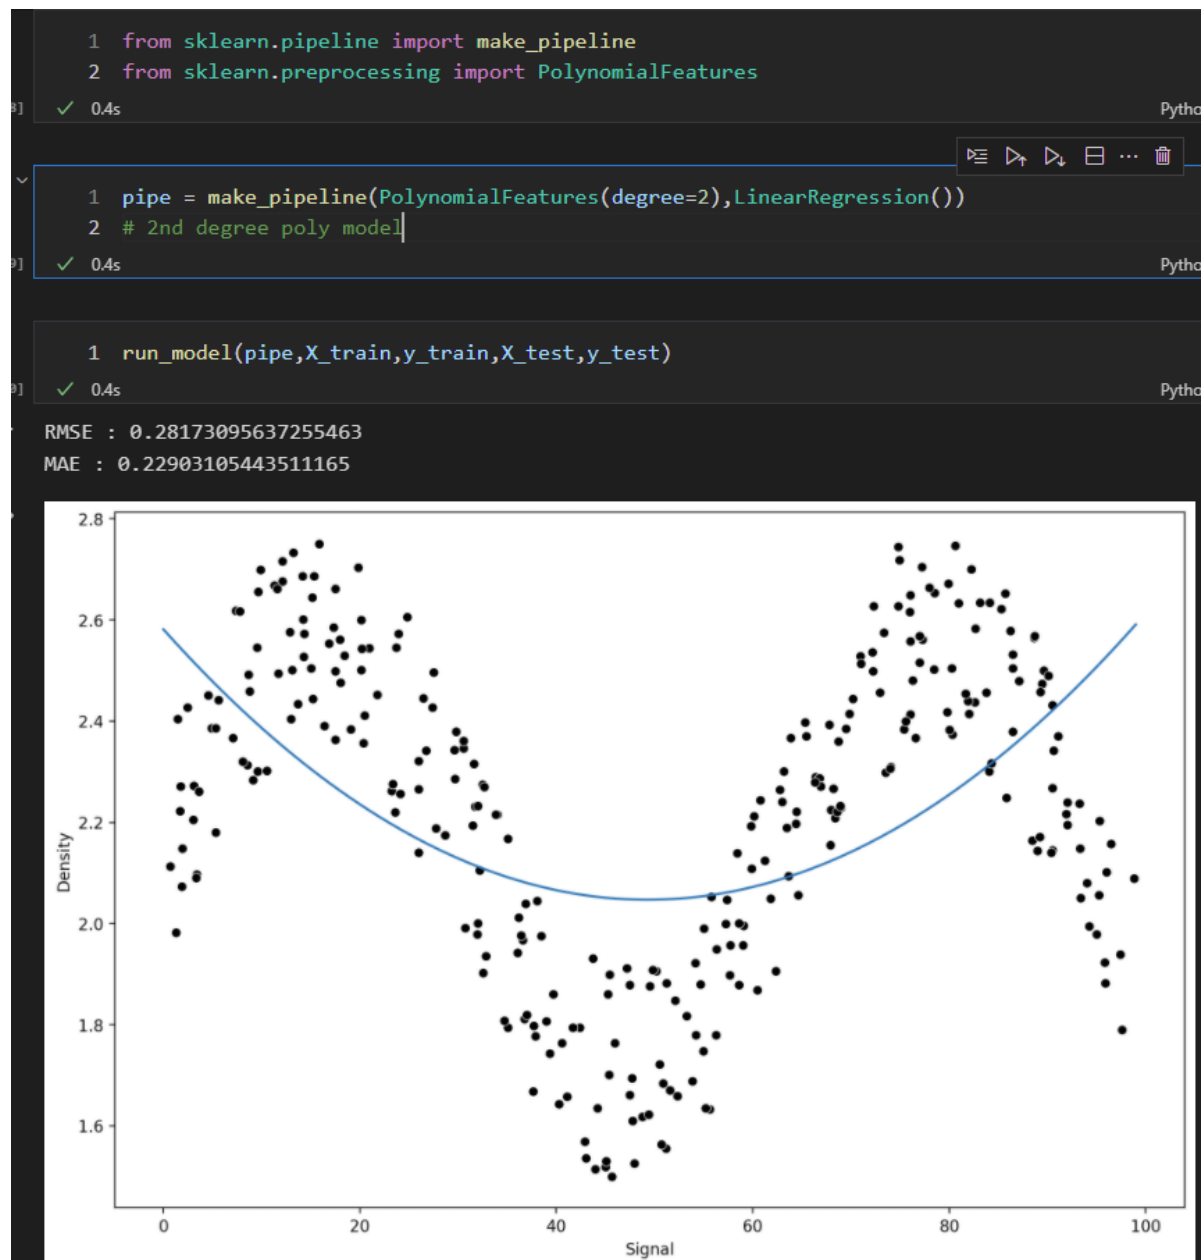
✓ 0.4s

Python

RMSE : 0.2570051996584629

MAE : 0.211198973318633






```
1 pipe = make_pipeline(PolynomialFeatures(degree=10),LinearRegression())  
2 run_model(pipe,X_train,y_train,X_test,y_test)
```

✓ 0.5s

Python

RMSE : 0.14034843686496898

MAE : 0.12467013990225653




```
1 pipe = make_pipeline(PolynomialFeatures(degree=4),LinearRegression())  
2 run_model(pipe,X_train,y_train,X_test,y_test)
```

✓ 0.6s

Pytho

RMSE : 0.1458863339756185

MAE : 0.1184764278843919



```
1 from sklearn.neighbors import KNeighborsRegressor
```

✓ 0.4s

```
1 k_values = [1,5,10,50]
```

```
2
```

```
3 for n in k_values:
```

```
4     model = KNeighborsRegressor(n_neighbors=n)
```

```
5     run_model(model,X_train,y_train,X_test,y_test)
```

✓ 2.3s

RMSE : 0.1523487028635337

MAE : 0.11877297474442378

RMSE : 0.13730685016923647

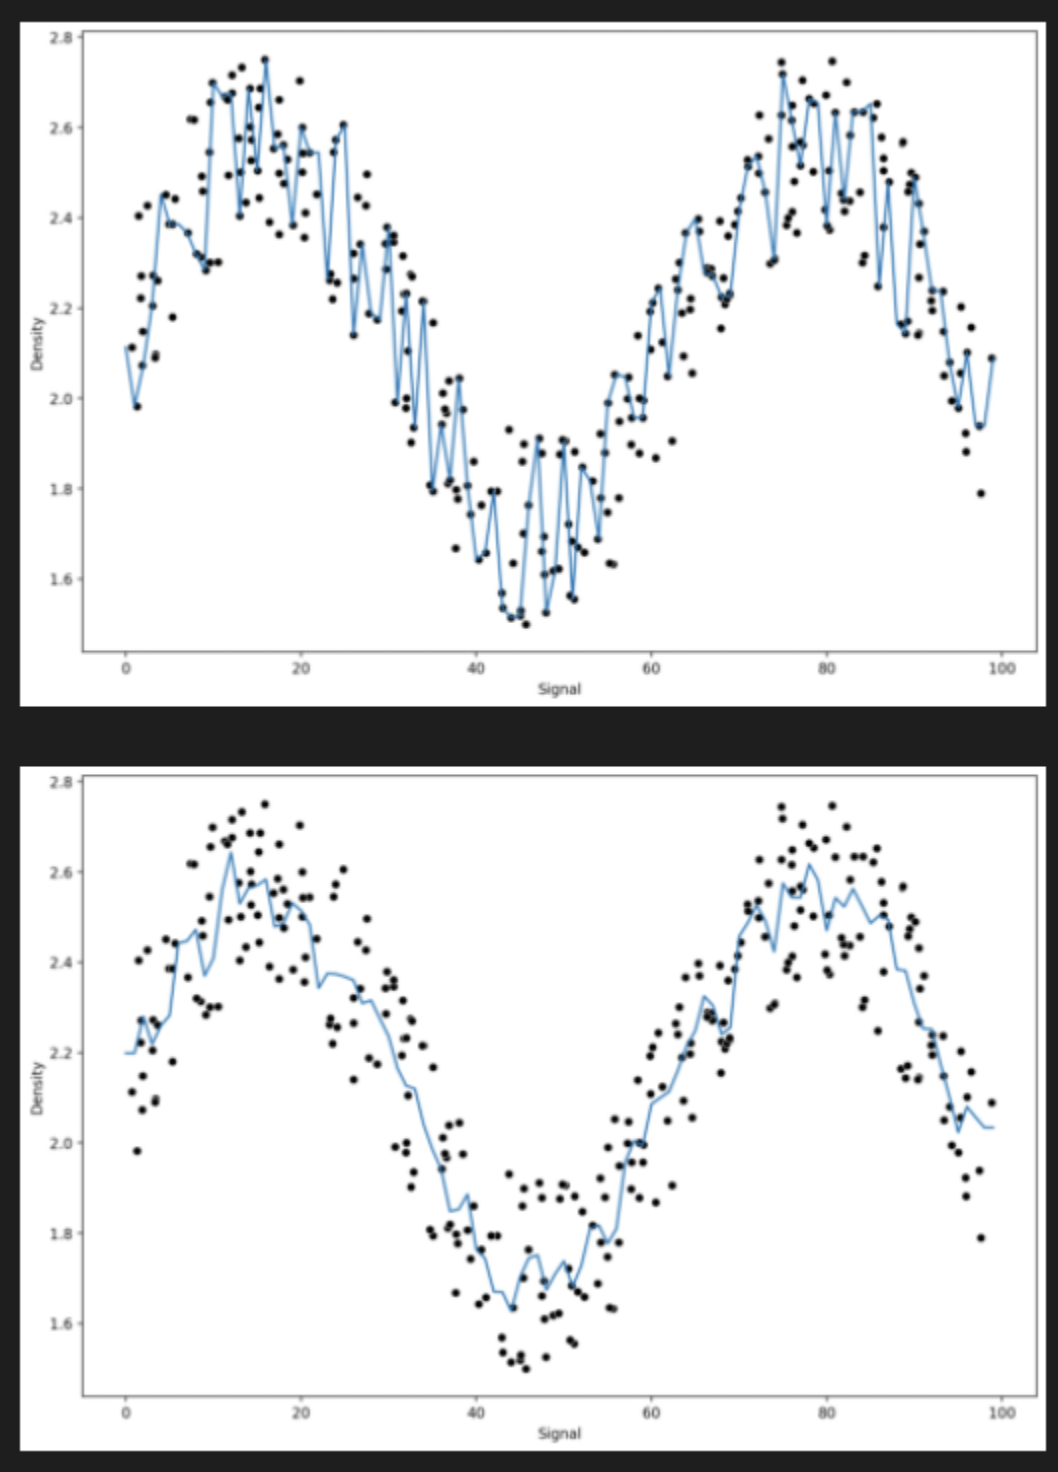
MAE : 0.12198383614100558

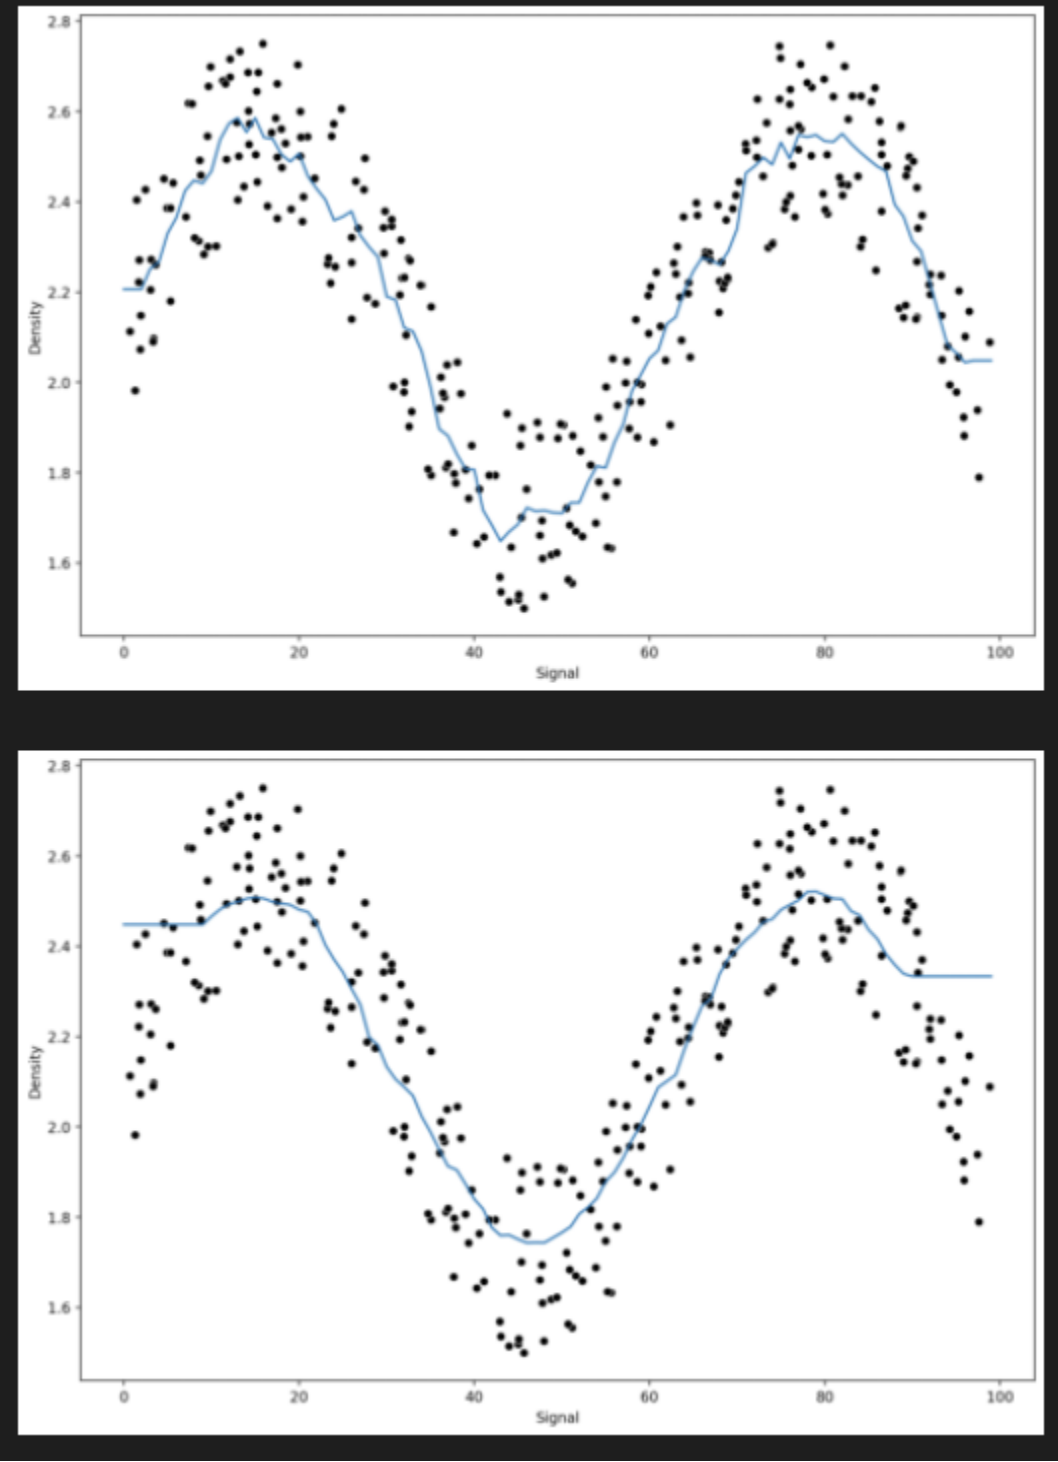
RMSE : 0.13277855732740926

MAE : 0.11635971693292672

RMSE : 0.19545005360281248

MAE : 0.1570937980156112







```
1 from sklearn.tree import DecisionTreeRegressor
```

✓ 0.3s

Pytho


```
1 model = DecisionTreeRegressor()  
2 run_model(model,X_train,y_train,X_test,y_test)
```

✓ 0.4s

Pytho

RMSE : 0.1523487028635337

MAE : 0.11877297474442378




```
1 from sklearn.svm import SVR #Support Vector Machines
2 from sklearn.model_selection import GridSearchCV
```

✓ 0.5s

+ Code

+ Markdown



```
1 svr = SVR()
2 param_grid = {
3     "C": [0.01, 0.1, 1, 5, 10, 100],
4     "gamma": ["auto", "scale"]
5 }
6 grid = GridSearchCV(svr, param_grid)
```

✓ 0.3s

Pytho


```
1 run_model(grid, X_train, y_train, X_test, y_test)
```

✓ 0.8s

Pytho

RMSE : 0.13015742723601528


MAE : 0.11243103141068958



```
1 from sklearn.ensemble import RandomForestRegressor
```

✓ 0.1s

Python




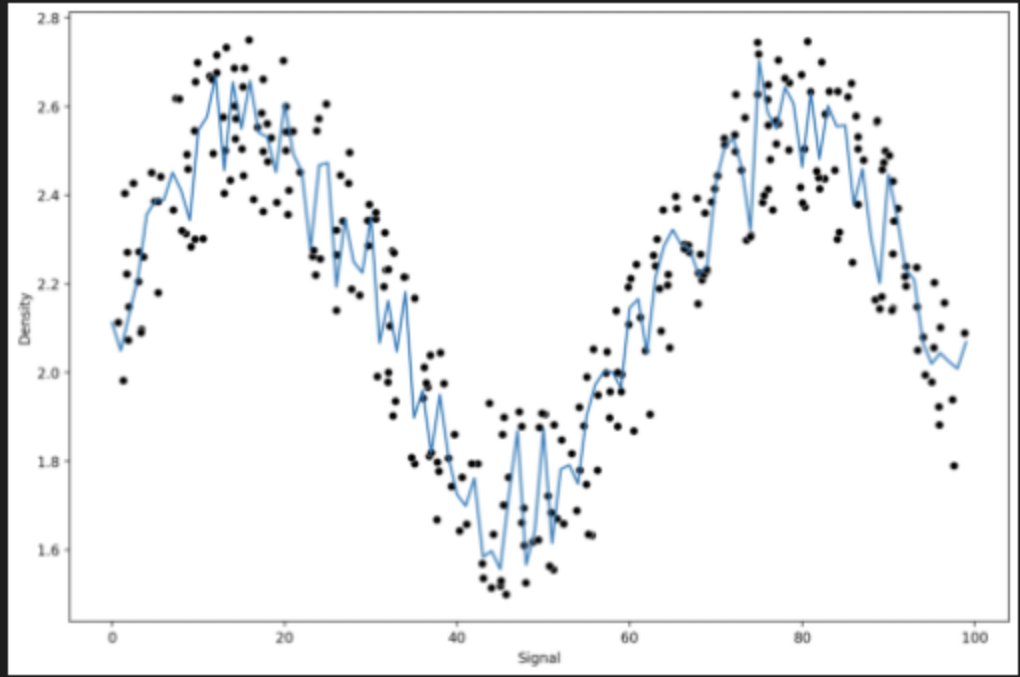

```
1 trees = [10,50,100]
2 for n in trees:
3     |
4     | model = RandomForestRegressor(n_estimators=n)
5     |
6     | run_model(model,X_train,y_train,X_test,y_test)
```

✓ 1.2s

Python

RMSE : 0.13385234470012897  
MAE : 0.10977912180641283  
RMSE : 0.13511144891461585  
MAE : 0.11339472588441249  
RMSE : 0.13372248344853374  
MAE : 0.11046317946299357





```
1 from sklearn.ensemble import GradientBoostingRegressor, AdaBoostRegressor
```

✓ 0.6s

Python


```
1 model = GradientBoostingRegressor()  
2 run_model(model,X_train,y_train,X_test,y_test)
```

✓ 0.5s

Python

RMSE : 0.13294148649584667

MAE : 0.11318284854800689



```
1 model = AdaBoostRegressor()  
2 run_model(model,X_train,y_train,X_test,y_test)
```

✓ 0.6s

Python

RMSE : 0.1361774372909145

MAE : 0.11587883158079351

