



Overview

Code

Markdown

TELCO CUSTOMER CHURN

Imports and Data

```
1 import pandas as pd
2 import numpy as np
3 import seaborn as sns
4 import matplotlib.pyplot as plt
```

✓ 1.8s

```
1 df = pd.read_csv("Telco-Customer-Churn.csv")
```

✓ 0.1s

```
1 df.head()
```

✓ 0.1s

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	In
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	
1	5575-GNVDE	Male	0	No	No	34	Yes	No	
2	3668-QPYBK	Male	0	No	No	2	Yes	No	

```
1 df.info()
```

✓ 0.1s

Output exceeds the [size limit](#). Open the full output data

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 7032 entries, 0 to 7031

Data columns (total 21 columns):

#	Column	Non-Null Count	Dtype
0	customerID	7032 non-null	object
1	gender	7032 non-null	object
2	SeniorCitizen	7032 non-null	int64
3	Partner	7032 non-null	object
4	Dependents	7032 non-null	object
5	tenure	7032 non-null	int64
6	PhoneService	7032 non-null	object
7	MultipleLines	7032 non-null	object
8	InternetService	7032 non-null	object
9	OnlineSecurity	7032 non-null	object
...			
19	TotalCharges	7032 non-null	float64
20	Churn	7032 non-null	object

dtypes: float64(2), int64(2), object(17)

memory usage: 1.1+ MB

```
1 df.isnull().sum().sum()
```

```
2 # Null değer yok
```

✓ 0.1s

0

```
1 df.isna().sum().sum()
```

```
2 # NaN değer yok
```

✓ 0.2s

0

```
1 df.describe()
```

✓ 0.1s


	SeniorCitizen	tenure	MonthlyCharges	TotalCharges
count	7032.000000	7032.000000	7032.000000	7032.000000
mean	0.162400	32.421786	64.798208	2283.300441
std	0.368844	24.545260	30.085974	2266.771362
min	0.000000	1.000000	18.250000	18.800000
25%	0.000000	9.000000	35.587500	401.450000
50%	0.000000	29.000000	70.350000	1397.475000
75%	0.000000	55.000000	89.862500	3794.737500
max	1.000000	72.000000	118.750000	8684.800000

EDA

```
1 sns.countplot(data=df, x="Churn")
```

✓ 0.3s

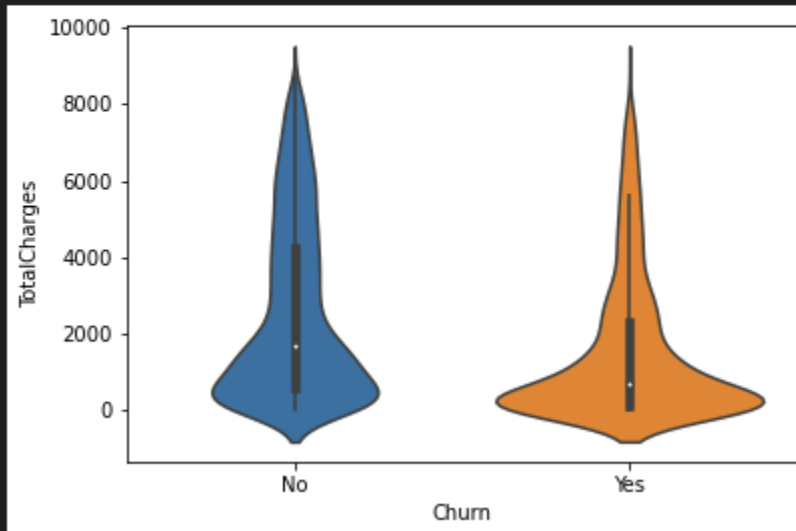
<AxesSubplot:xlabel='Churn', ylabel='count'>



```
1 sns.violinplot(data=df, x="Churn", y="TotalCharges")
```

✓ 0.4s

```
<AxesSubplot:xlabel='Churn', ylabel='TotalCharges'>
```




```
1 plt.figure(figsize=(10,4),dpi=200)
2 sns.boxplot(data=df, y="TotalCharges", x="Contract", hue="Churn")
3 plt.legend(loc=(1.1,0.5))
```

✓ 0.4s

Python

```
<matplotlib.legend.Legend at 0x22ad2c24e50>
```



```
1 df.columns
```

```
✓ 0.7s
```

```
Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',  
      'tenure', 'PhoneService', 'MultipleLines', 'InternetService',  
      'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',  
      'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',  
      'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],  
      dtype='object')
```

```
1 corr = pd.get_dummies(df[['gender', 'SeniorCitizen', 'Partner',  
2   'Dependents', 'PhoneService', 'MultipleLines', 'InternetService',  
3   'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',  
4   'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',  
5   'PaymentMethod', 'Churn']]).corr()
```

```
✓ 0.1s
```

```

1 corr["Churn_Yes"].sort_values().iloc[1:-1]
2 # Churn_Yes kendisi ile 1 tam korale Churn_No ile de1 tam korale
3 # ... oolduđu için 1. ve sonundu indeksteki veriler bizim için anlam
4 # ... ifade etmez. Bunları atmak için de iloc[1:-1] kullandık.
5 # ... yani 0. indeksi dışarıda bırakacak şekilde 1den -1. indekse kadar
6 # ... olan serinin tamamını aldık.

```

✓ 0.1s

Output exceeds the [size limit](#). Open the full output data [in a text editor](#)

Contract_Two year	-0.301552
StreamingMovies_No internet service	-0.227578
StreamingTV_No internet service	-0.227578
TechSupport_No internet service	-0.227578
DeviceProtection_No internet service	-0.227578
OnlineBackup_No internet service	-0.227578
OnlineSecurity_No internet service	-0.227578
InternetService_No	-0.227578
PaperlessBilling_No	-0.191454
Contract_One year	-0.178225
OnlineSecurity_Yes	-0.171270
TechSupport_Yes	-0.164716
Dependents_Yes	-0.163128
Partner_Yes	-0.149982
PaymentMethod_Credit card (automatic)	-0.134687
...	
InternetService_Fiber optic	0.307463
TechSupport_No	0.336877
OnlineSecurity_No	0.342235
Contract_Month-to-month	0.404565

Name: Churn_Yes, dtype: float64


```

1 plt.figure(figsize=(10,4),dpi=200)
2 sns.barplot(x=corr["Churn_Yes"].sort_values().iloc[1:-1].index,
3 | y=corr['Churn_Yes'].sort_values().iloc[1:-1].values)
4 plt.title("Feature Correlation to Yes Churn")
5 plt.xticks(rotation=90);

```

✓ 1.5s

Python



Churn Analysis

```
1 df["Contract"].unique()
```

✓ 0.7s

Python

```
array(['Month-to-month', 'One year', 'Two year'], dtype=object)
```

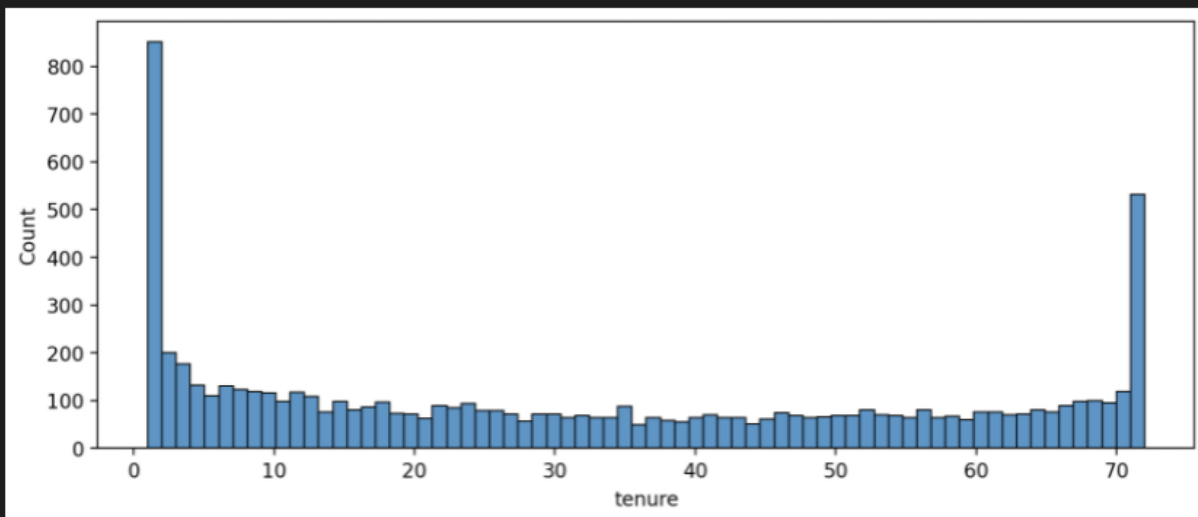
```
1 plt.figure(figsize=(10,4),dpi=200)
```

```
2 sns.histplot(data=df, x="tenure", bins=70)
```

✓ 0.6s

Python

```
<AxesSubplot:xlabel='tenure', ylabel='Count'>
```

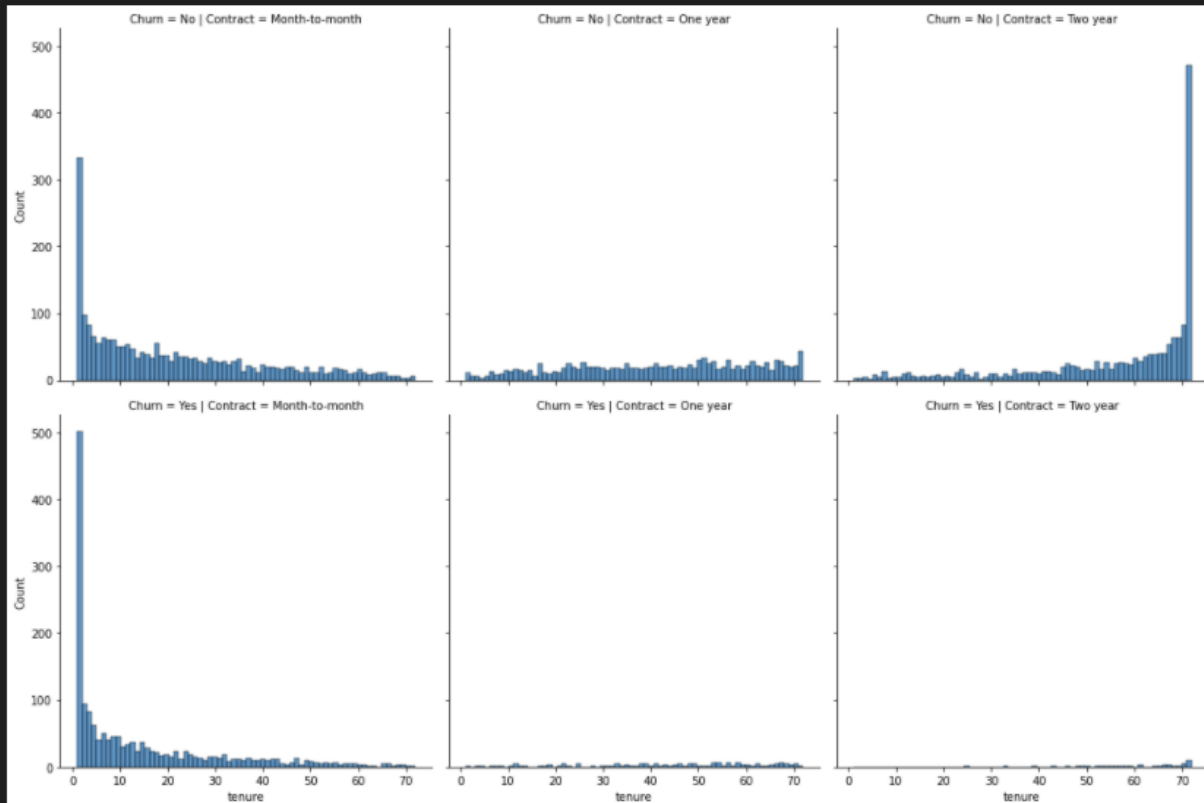



```
1 plt.figure(figsize=(10,4),dpi=200)
2 sns.displot(data=df, x="tenure",bins=70, col="Contract", row="Churn");
```

✓ 2.7s

Python

<Figure size 2000x800 with 0 Axes>




```

1 plt.figure(figsize=(10,4),dpi=200)
2 sns.scatterplot(data=df, x="MonthlyCharges", y="TotalCharges",
3 | hue="Churn", lw=0.5, alpha=0.6);

```

✓ 0.7s

Python



```

1 no_churn = df.groupby(["Churn","tenure"]).count().transpose()["No"]
2 yes_churn = df.groupby(["Churn","tenure"]).count().transpose()["Yes"]

```

✓ 0.1s

Python

```

1 yes_churn

```

✓ 0.1s

Python

	tenure	1	2	3	4	5	6	7	8	9	10	...	63	64	65	66	67	68	69	70	71
customerID	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8	11	6	
gender	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8	11	6	
SeniorCitizen	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8	11	6	
Partner	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8	11	6	
Dependents	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8	11	6	
PhoneService	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8	11	6	
MultipleLines	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8	11	6	
InternetService	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8	11	6	
OnlineSecurity	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8	11	6	
OnlineBackup	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8	11	6	

```
1 churn_rate = 100 * yes_churn / (yes_churn + no_churn)
```

✓ 0.9s

```
1 churn_rate.transpose()["customerID"]
```

✓ 0.1s

tenure

1	61.990212
2	51.680672
3	47.000000
4	47.159091
5	48.120301

...

68	9.000000
69	8.421053
70	9.243697
71	3.529412
72	1.657459


Name: customerID, Length: 72, dtype: float64

```
1 plt.figure(figsize=(10,4),dpi=200)
2 churn_rate.iloc[0].plot()
3 plt.ylabel("Churn Percentage")
```

✓ 0.3s

Python


Text(0, 0.5, 'Churn Percentage')



```
1 churn_rate.transpose()["customerID"].plot()
```

✓ 0.3s

<AxesSubplot:xlabel='tenure'>



```

1 def cohort(tenure):
2     if tenure < 13:
3         return "0-12 Months"
4     elif tenure < 25:
5         return '12-24 Months'
6     elif tenure < 49:
7         return '24-48 Months'
8     else:
9         return "Over 48 Months"

```

5] ✓ 0.1s

```

1 df['Tenure Cohort'] = df['tenure'].apply(cohort)

```

5] ✓ 0.2s

```

1 df.head(10)[['tenure', 'Tenure Cohort']]

```

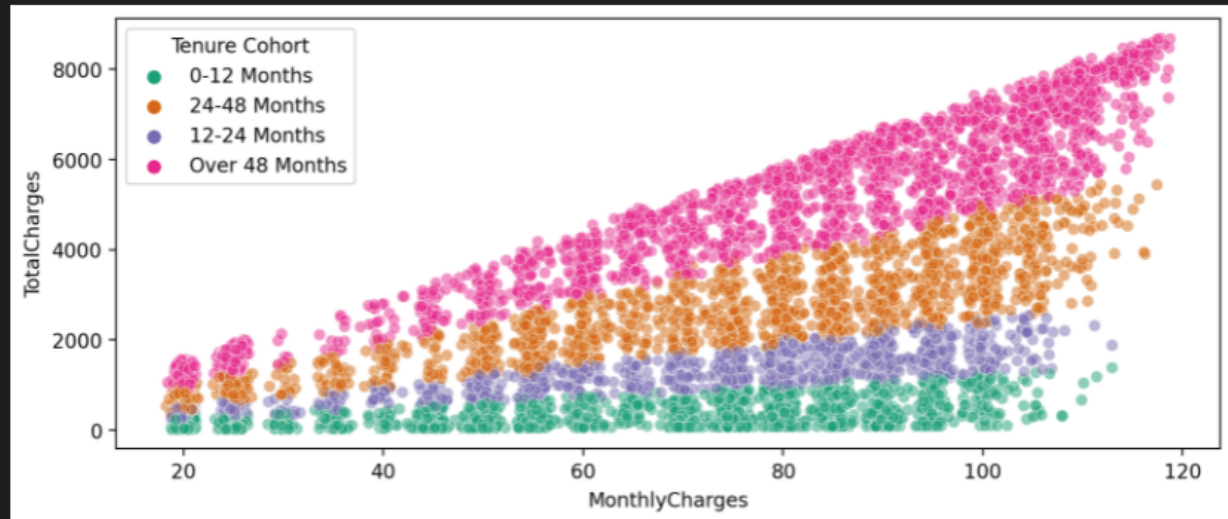
7] ✓ 0.1s

	tenure	Tenure Cohort
0	1	0-12 Months
1	34	24-48 Months
2	2	0-12 Months
3	45	24-48 Months
4	2	0-12 Months
5	8	0-12 Months
6	22	12-24 Months
7	10	0-12 Months
8	28	24-48 Months
9	62	Over 48 Months

```
1 plt.figure(figsize=(10,4),dpi=200)
2 sns.scatterplot(data=df,x='MonthlyCharges',y='TotalCharges',
3               hue='Tenure Cohort', linewidth=0.5,alpha=0.5,palette='Dark2')
✓ 1.2s
```

Python

<AxesSubplot:xlabel='MonthlyCharges', ylabel='TotalCharges'>

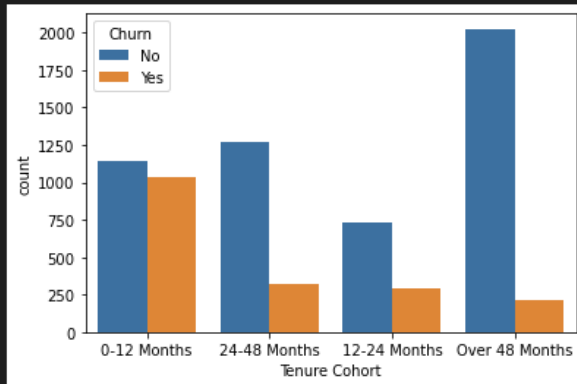


```
1 sns.countplot(data=df, x="Tenure Cohort", hue="Churn")
```

✓ 0.7s

Python

<AxesSubplot:xlabel='Tenure Cohort', ylabel='count'>




```
1 sns.catplot(data=df, x="Tenure Cohort", hue="Churn",  
2 | kind="count", col="Contract")
```

✓ 0.8s

Python

<seaborn.axisgrid.FacetGrid at 0x22ad3a98d00>



Single Decision Tree

```
1 X = df.drop(["Churn","customerID"], axis=1)
2 y = df["Churn"]
```

✓ 0.6s

Python

```
1 X = pd.get_dummies(X, drop_first=True)
```

✓ 0.1s

Python

```
1 from sklearn.model_selection import train_test_split
```

✓ 0.3s

Python

```
1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, random_state=101)
```

✓ 0.1s

Python

```
1 from sklearn.tree import DecisionTreeClassifier
```

✓ 0.1s

Python

```
1 dt = DecisionTreeClassifier(max_depth=6)
```

✓ 0.9s

Python

```
1 dt.fit(X_train, y_train)
```

✓ 0.2s

Python

```
DecisionTreeClassifier(max_depth=6)
```

```
1 preds = dt.predict(X_test)
```

✓ 0.1s

Python


```
1 from sklearn.metrics import classification_report, plot_confusion_matrix, accuracy_score
```

✓ 0.7s

```
1 print(classification_report(y_test,preds))
```


✓ 0.1s

	precision	recall	f1-score	support
No	0.87	0.89	0.88	557
Yes	0.55	0.49	0.52	147
accuracy			0.81	704
macro avg	0.71	0.69	0.70	704
weighted avg	0.80	0.81	0.81	704

```
1 plot_confusion_matrix(dt,X_test,y_test)
```

✓ 0.8s

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x22ad2e4d160>



```

1 imp_feats = pd.DataFrame([data=dt.feature_importances_, index=X.columns,
2 | columns=["Feature Importance"]]).sort_values("Feature Importance")

```

✓ 0.1s

Python


```

1 plt.figure(figsize=(14,6),dpi=200)
2 sns.barplot(data=imp_feats.sort_values('Feature Importance'),
3 | x=imp_feats.sort_values('Feature Importance').index,y='Feature Importance')
4 plt.xticks(rotation=90)
5 plt.title("Feature Importance for Decision Tree");

```

✓ 1.1s

Python



```
1 from sklearn.tree import plot_tree
```


✓ 0.2s

Python

```
1 plt.figure(figsize=(12,8),dpi=150)  
2 plot_tree(dt,filled=True, feature_names=X.columns);
```

✓ 94s

Python



Random Forest

```
1 from sklearn.ensemble import RandomForestClassifier
```

✓ 0.9s

```
1 rf = RandomForestClassifier(n_estimators=100)
```

✓ 0.1s

```
1 rf.fit(X_train, y_train)
```

✓ 1.3s

RandomForestClassifier()

```
1 preds = rf.predict(X_test)
```

✓ 0.1s

```
1 print(classification_report(y_test, preds))
```


✓ 0.1s

	precision	recall	f1-score	support
No	0.86	0.89	0.87	557
Yes	0.52	0.46	0.49	147
accuracy			0.80	704
macro avg	0.69	0.67	0.68	704
weighted avg	0.79	0.80	0.79	704

```
1 plot_confusion_matrix(rf,X_test,y_test)
```

✓ 0.6s

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatri
```



Boosted Trees

```
1 from sklearn.ensemble import GradientBoostingClassifier, AdaBoostClassifier
```

✓ 0.9s

```
1 ada_model = AdaBoostClassifier()
```

✓ 0.1s

```
1 ada_model.fit(X_train, y_train)
```

✓ 1.3s

AdaBoostClassifier()

```
1 preds = ada_model.predict(X_test)
```

✓ 0.1s

```
1 print(classification_report(y_test, preds))
```

✓ 0.2s

	precision	recall	f1-score	support
No	0.88	0.90	0.89	557
Yes	0.60	0.54	0.57	147
accuracy			0.83	704
macro avg	0.74	0.72	0.73	704
weighted avg	0.82	0.83	0.83	704

```
1 plot_confusion_matrix(rf, X_test, y_test)
```

✓ 0.6s

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMa
```

