



Hierarchical Clustering

▼ Clustering

Hierarchical Clustering

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
```

1] ✓ 9.8s Python

```
1 df = pd.read_csv('cluster_mpg.csv')
```

2] ✓ 0.6s Python

```
1 df
```

3] ✓ 0.6s Python

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin	name
0	18.0	8	307.0	130.0	3504	12.0	70	usa	chevrolet chevelle malibu
1	15.0	8	350.0	165.0	3693	11.5	70	usa	buick skylark 320
2	18.0	8	318.0	150.0	3436	11.0	70	usa	plymouth satellite
3	16.0	8	304.0	150.0	3433	12.0	70	usa	amc rebel sst
4	17.0	8	302.0	140.0	3449	10.5	70	usa	ford torino
...

Data & EDA

```
1 df.describe()
```

0.2s

Python

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year
count	392.000000	392.000000	392.000000	392.000000	392.000000	392.000000	392.000000
mean	23.445918	5.471939	194.411990	104.469388	2977.584184	15.541327	75.979592
std	7.805007	1.705783	104.644004	38.491160	849.402560	2.758864	3.683737
min	9.000000	3.000000	68.000000	46.000000	1613.000000	8.000000	70.000000
25%	17.000000	4.000000	105.000000	75.000000	2225.250000	13.775000	73.000000
50%	22.750000	4.000000	151.000000	93.500000	2803.500000	15.500000	76.000000
75%	29.000000	8.000000	275.750000	126.000000	3614.750000	17.025000	79.000000
max	46.600000	8.000000	455.000000	230.000000	5140.000000	24.800000	82.000000

```
1 df["origin"].value_counts()
```

0.2s

Python

```
usa      245
japan    79
europe   68
Name: origin, dtype: int64
```

```
1 df_w_dummies = pd.get_dummies(df.drop("name", axis=1))
```

0.1s

Python

1 df_w_dummies

✓ 0.2s Python

s	displacement	horsepower	weight	acceleration	model_year	origin_europe	origin_japan
8	307.0	130.0	3504	12.0	70	0	0
8	350.0	165.0	3693	11.5	70	0	0
8	318.0	150.0	3436	11.0	70	0	0
8	304.0	150.0	3433	12.0	70	0	0
8	302.0	140.0	3449	10.5	70	0	0
..
4	140.0	86.0	2790	15.6	82	0	0
4	97.0	52.0	2130	24.6	82	1	0
4	135.0	84.0	2295	11.6	82	0	0
4	120.0	79.0	2625	18.6	82	0	0
4	119.0	82.0	2720	19.4	82	0	0

s

```
1 from sklearn.preprocessing import MinMaxScaler
```

[8] ✓ 1.4s

```
1 scaler = MinMaxScaler()
```

[9] ✓ 0.2s

```
1 scaled_data = scaler.fit_transform(df_w_dummies)
```

[10] ✓ 0.1s

```
1 scaled_data
```

[11] ✓ 0.1s

```
.. array([[0.2393617 , 1.          , 0.61757106, ..., 0.          , 0.          ,
          1.          ],
          [0.15957447, 1.          , 0.72868217, ..., 0.          , 0.          ,
          1.          ],
          [0.2393617 , 1.          , 0.64599483, ..., 0.          , 0.          ,
          1.          ],
          ...,
          [0.61170213, 0.2          , 0.17312661, ..., 0.          , 0.          ,
          1.          ],
          [0.50531915, 0.2          , 0.13436693, ..., 0.          , 0.          ,
          1.          ],
          [0.58510638, 0.2          , 0.13178295, ..., 0.          , 0.          ,
          1.          ]])
```

```
1 scaled_df = pd.DataFrame(scaled_data, columns=df_w_dummies.columns)
```

✓ 0.8s Python

```
1 scaled_df
```

✓ 0.1s Python

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	c
0	0.239362	1.0	0.617571	0.456522	0.536150	0.238095	0.0	
1	0.159574	1.0	0.728682	0.646739	0.589736	0.208333	0.0	
2	0.239362	1.0	0.645995	0.565217	0.516870	0.178571	0.0	
3	0.186170	1.0	0.609819	0.565217	0.516019	0.238095	0.0	
4	0.212766	1.0	0.604651	0.510870	0.520556	0.148810	0.0	
...
387	0.478723	0.2	0.186047	0.217391	0.333711	0.452381	1.0	
388	0.930851	0.2	0.074935	0.032609	0.146583	0.988095	1.0	
389	0.611702	0.2	0.173127	0.206522	0.193365	0.214286	1.0	
390	0.505319	0.2	0.134367	0.179348	0.286929	0.630952	1.0	
391	0.585106	0.2	0.131783	0.195652	0.313864	0.678571	1.0	


392 rows × 10 columns

```
1 plt.figure(figsize=(12,8), dpi=200)
2 sns.heatmap(scaled_df, cmap="viridis")
```

✓ 2.4s

Python

<AxesSubplot:>




```
1 plt.figure(figsize=(12,12), dpi=200)
2 sns.clustermap(scaled_df, cmap="viridis")
```

✓ 1.7s

Pyth


<seaborn.matrix.ClusterGrid at 0x207bc4f6970>

<Figure size 2400x2400 with 0 Axes>



1
scaled_df.corr()
0.1s
Python

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin_europe	origin_japan	origin_usa
mpg	1.000000	-0.777618	-0.805127	-0.778427	-0.832244	0.423329	0.580541	0.244313	0.451454	-0.565161
cylinders	-0.777618	1.000000	0.950823	0.842983	0.897527	-0.504683	-0.345647	-0.352324	-0.404209	0.610494
displacement	-0.805127	0.950823	1.000000	0.897257	0.932994	-0.543800	-0.369855	-0.371633	-0.440825	0.655936
horsepower	-0.778427	0.842983	0.897257	1.000000	0.864538	-0.689196	-0.416361	-0.284948	-0.321936	0.489625
weight	-0.832244	0.897527	0.932994	0.864538	1.000000	-0.416839	-0.309120	-0.293841	-0.447929	0.600978
acceleration	0.423329	-0.504683	-0.543800	-0.689196	-0.416839	1.000000	0.290316	0.208298	0.115020	-0.258224
model_year	0.580541	-0.345647	-0.369855	-0.416361	-0.309120	0.290316	1.000000	0.000000	0.000000	0.000000
origin_europe	0.244313	-0.352324	-0.371633	-0.284948	-0.293841	0.208298	0.000000	1.000000	0.000000	0.000000
origin_japan	0.451454	-0.404209	-0.440825	-0.321936	-0.447929	0.115020	0.000000	0.000000	1.000000	0.000000
origin_usa	-0.565161	0.610494	0.655936	0.489625	0.600978	-0.258224	0.000000	0.000000	0.000000	1.000000



```

1 plt.figure(figsize=(12,12), dpi=200)
2 sns.clustermap(scaled_df, cmap="viridis", col_cluster=False)


```

✓ 2.5s

Pyth

<seaborn.matrix.ClusterGrid at 0x207b99d29d0>

<Figure size 2400x2400 with 0 Axes>



ML Model

```
1 from sklearn.cluster import AgglomerativeClustering
```

✓ 0.4s

```
1 model = AgglomerativeClustering(n_clusters=4)
```

✓ 0.7s

```
1 cluster_labels = model.fit_predict(scaled_df)
```

✓ 0.1s

```
1 cluster_labels
```

✓ 0.9s


```
array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 3, 0, 0, 0, 3, 2, 2, 2,
       2, 2, 0, 1, 1, 1, 1, 3, 0, 3, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1,
       0, 0, 0, 0, 0, 2, 2, 2, 3, 3, 2, 0, 3, 0, 2, 0, 0, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 3, 1, 1, 1, 1, 2, 2, 2, 2, 0, 3, 3, 0, 3, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 2, 1, 1, 1, 1, 0, 3, 0, 3,
       3, 0, 0, 2, 1, 1, 2, 2, 2, 2, 1, 2, 3, 1, 0, 0, 0, 3, 0, 3, 0, 0,
       0, 0, 1, 1, 1, 1, 1, 2, 2, 2, 3, 3, 0, 2, 2, 3, 3, 2, 0, 0, 0, 0,
       1, 1, 1, 1, 0, 0, 0, 0, 0, 1, 1, 3, 0, 0, 0, 3, 2, 3, 0, 2, 0, 2,
       2, 2, 2, 3, 2, 2, 0, 0, 2, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 2, 3, 0,
       0, 0, 0, 2, 3, 3, 0, 2, 1, 2, 3, 2, 1, 1, 1, 1, 3, 0, 2, 0, 3, 1,
       1, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 2, 0, 3, 0, 0, 0, 3, 2, 3, 2, 3,
       2, 0, 3, 3, 3, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1,
```

```
1 plt.figure(figsize=(12,4),dpi=200)
2 sns.scatterplot(data=df, x="mpg", y="weight", hue=cluster_labels)
```

✓ 0.9s

Python

<AxesSubplot:xlabel='mpg', ylabel='weight'>



Dendrograms

documentation

<https://docs.scipy.org/doc/scipy/reference/generated/scipy.cluster.hierarchy.dendro>

```
1 model = AgglomerativeClustering(n_clusters=None, distance_threshold=0)
```

✓ 0.1s

```
1 cluster_labels = model.fit_predict(scaled_df)
```

✓ 0.1s

```
1 cluster_labels
```

✓ 0.1s

Output exceeds the [size limit](#). Open the full output data [in a text editor](#)

```
array([247, 252, 360, 302, 326, 381, 384, 338, 300, 279, 217, 311, 377,
       281, 232, 334, 272, 375, 354, 333, 317, 345, 329, 289, 305, 383,
       290, 205, 355, 269, 202, 144, 245, 297, 386, 358, 199, 337, 330,
       339, 293, 352, 283, 196, 253, 168, 378, 331, 201, 268, 256, 361,
       250, 197, 246, 371, 324, 230, 203, 261, 380, 376, 308, 389, 332,
       306, 236, 391, 350, 274, 288, 313, 231, 298, 100, 295, 210, 248,
       187, 390, 373, 266, 307, 379, 212, 357, 191, 314, 208, 249, 343,
       294, 374, 322, 323, 362, 188, 296, 369, 286, 251, 229, 244, 285,
       349, 365, 259, 213, 276, 215, 222, 204, 359, 287, 166, 387, 291,
       220, 216, 260, 129, 367, 340, 346, 301, 342, 228, 388, 370, 218,
       255, 327, 347, 278, 271, 258, 282, 318, 273, 123, 172, 382, 363,
       356, 195, 280, 239, 364, 267, 351, 186, 257, 277, 299, 127, 366,
       234, 385, 192, 372, 292, 233, 270, 263, 133, 165, 161, 198, 97,
       315, 134, 207, 147, 175, 262, 348, 98, 214, 48, 353, 177, 325,
       128, 284, 275, 182, 184, 145, 344, 321, 200, 149, 240, 241, 235,
```

...

```
1 from scipy.cluster.hierarchy import dendrogram
2 from scipy.cluster import hierarchy

27] ✓ 0.1s

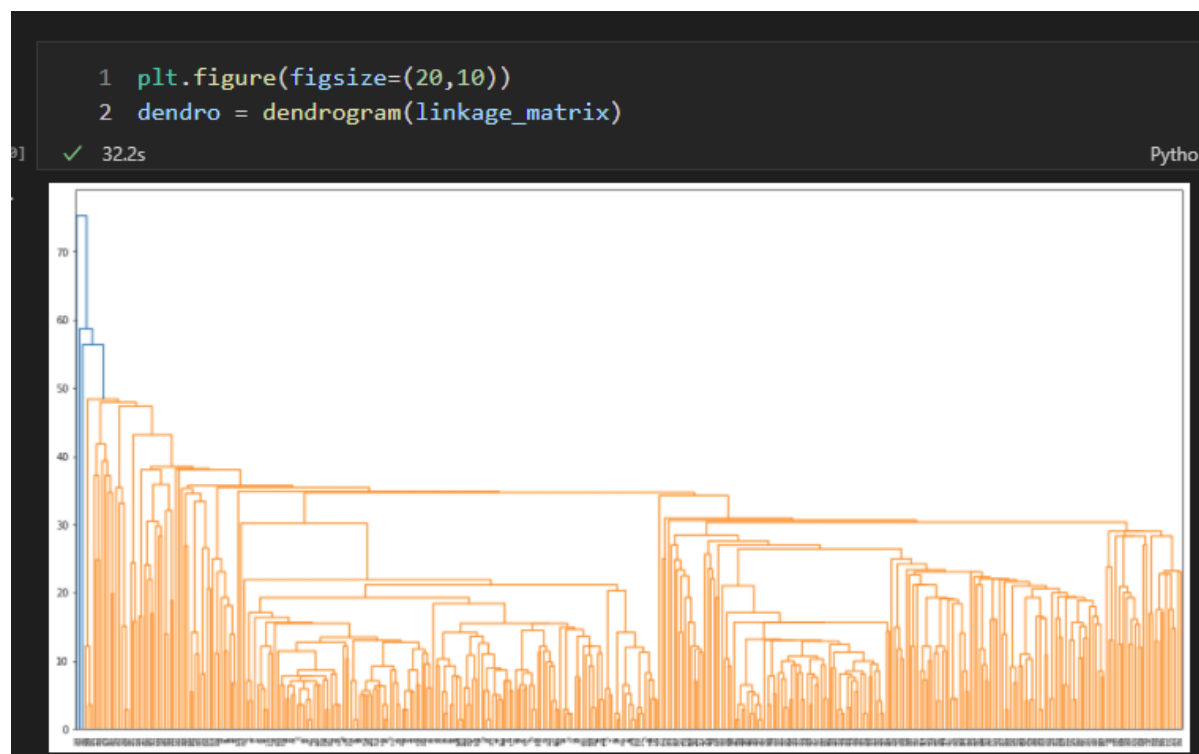
1 linkage_matrix = hierarchy.linkage(model.children_)

28] ✓ 0.1s

1 linkage_matrix

29] ✓ 0.1s


array([[ 67.      , 161.      ,  1.41421356,  2.      ],
       [ 10.      ,  45.      ,  1.41421356,  2.      ],
       [ 47.      ,  99.      ,  1.41421356,  2.      ],
       ...,
       [340.      , 777.      , 56.40035461, 389.      ],
       [332.      , 778.      , 58.69412236, 390.      ],
       [349.      , 779.      , 75.32595834, 391.      ]])
```



```
1 plt.figure(figsize=(20,10))  
2 dendro = dendrogram(linkage_matrix, truncate_mode="lastp", p=10)
```

✓ 0.3s

Python



Threshold Choosing

[+ Code](#)[+ Markdown](#)

```
1 scaled_df.describe()
```

✓ 0.1s

	mpg	cylinders	displacement	horsepower	weight	acceleration
count	392.000000	392.000000	392.000000	392.000000	392.000000	392.000000
mean	0.384200	0.494388	0.326646	0.317768	0.386897	0.386897
std	0.207580	0.341157	0.270398	0.209191	0.240829	0.240829
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.212766	0.200000	0.095607	0.157609	0.173589	0.173589
50%	0.365691	0.200000	0.214470	0.258152	0.337539	0.337539
75%	0.531915	1.000000	0.536822	0.434783	0.567550	0.567550
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

```
1 scaled_df["mpg"].idxmax()
```

✓ 0.7s

320

```
1 scaled_df["mpg"].idxmin()
```

✓ 0.1s

28


```
1 # https://stackoverflow.com/questions/1401712/how-can-tl
2
3 a = scaled_df.iloc[320]
4 b = scaled_df.iloc[28]
5
✓ 0.1s
```

```
1 np.sqrt(len(scaled_df.columns))
✓ 0.1s
```

3.1622776601683795

```
1 distance = np.linalg.norm(a-b)
✓ 0.1s
```

```
1 distance
✓ 0.8s
```

2.3852929970374714

```
1 model = AgglomerativeClustering(n_clusters=None,
2 | distance_threshold=2)
✓ 0.1s Python
```

```
1 cluster_labels = model.fit_predict(scaled_data)
✓ 0.1s Python
```

```
1 cluster_labels
✓ 0.1s Python
```

Output exceeds the [size limit](#). Open the full output data [in a text editor](#)

```
array([ 3,  3,  3,  3,  3,  3,  3,  3,  3,  3,  3,  3,  3,  3,  1,  4,
4,
         4,  1,  0,  0,  0,  0,  0,  4,  3,  3,  3,  3,  1,  7,  1,  4,
4,
         4,  4,  4,  3,  3,  3,  3,  3,  3,  3,  3,  4,  7,  4,  4,  7,  0,
0,
         0,  1,  1,  0,  7,  1,  7,  0,  7,  7,  3,  3,  3,  3,  3,  3,
3,
         3,  3,  1,  3,  3,  3,  3,  0,  0,  0,  0,  7,  1,  1,  7,  1,
3,
         3,  3,  3,  3,  3,  3,  3,  3,  3,  3,  3,  4,  4,  4,  4,  4,
0,
         3,  3,  3,  3,  4,  1,  7,  1,  1,  7,  4,  0,  3,  3,  0,  0,
0,
```

```
1 np.unique(cluster_labels)
✓ 0.8s Python
```

```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10], dtype=int64)
```

Linkage Matrix

Source:

<https://docs.scipy.org/doc/scipy/reference/generated/scipy.cluster.hierarchy.>

A $(n-1)$ by 4 matrix Z is returned. At the i -th iteration, clusters with indices $Z[i, 0]$ and $Z[i, 1]$ are combined to form cluster $n + i$. A cluster with an index less than n corresponds to one of the original observations. The distance between clusters $Z[i, 0]$ and $Z[i, 1]$ is given by $Z[i, 2]$. The fourth value $Z[i, 3]$ represents the number of original observations in the newly formed cluster.

```
1 linkage_matrix = hierarchy.linkage(model.children_)
```

✓ 0.9s

Python

```
1 linkage_matrix
```

✓ 0.9s

Python

```
array([[ 67.      , 161.      ,  1.41421356,  2.      ],
       [ 10.      ,  45.      ,  1.41421356,  2.      ],
       [ 47.      ,  99.      ,  1.41421356,  2.      ],
       ...,
       [340.      , 777.      , 56.40035461, 389.      ],
       [332.      , 778.      , 58.69412236, 390.      ],
       [349.      , 779.      , 75.32595834, 391.      ]])
```

