



DBSCAN

▼ DBSCAN vs K-Means

DBSCAN and Clustering Examples

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
```

✓ 10.5s

```
1 blobs = pd.read_csv('cluster_blobs.csv')
```

✓ 0.9s

```
1 blobs.head()
```


✓ 0.2s

	X1	X2
0	4.645333	6.822294
1	4.784032	6.422883
2	-5.851786	5.774331
3	-7.459592	6.456415
4	4.918911	6.961479

```
1 sns.scatterplot(data=blobs, x="X1", y="X2")
```

✓ 0.7s

```
<AxesSubplot:xlabel='X1', ylabel='X2'>
```



```
1 moons = pd.read_csv('cluster_moons.csv')
```

✓ 0.9s

```
1 moons.tail()
```


✓ 0.1s

	X1	X2
1495	1.957344	0.187184
1496	0.962394	0.384304
1497	-0.761893	0.581666
1498	1.803858	-0.154705
1499	0.203305	0.079049

```
1 sns.scatterplot(data=moons, x="X1", y="X2")
```

✓ 0.6s

<AxesSubplot:xlabel='X1', ylabel='X2'>



```
1 circles = pd.read_csv('cluster_circles.csv')
```

✓ 0.9s

```
1 circles.head()
```


✓ 0.1s

	X1	X2
0	-0.348677	0.010157
1	-0.176587	-0.954283
2	0.301703	-0.113045
3	-0.782889	-0.719468
4	-0.733280	-0.757354

```
1 sns.scatterplot(data=circles, x="X1", y="X2")
```

✓ 0.6s

<AxesSubplot:xlabel='X1', ylabel='X2'>



```

1 def display_categories(model,data):
2
3     labels = model.fit_predict(data)
4     sns.scatterplot(data=data, x='X1', y='X2', hue=labels, palette="Set1")

```

✓ 0.1s

K_Means

```

1 from sklearn.cluster import KMeans

```


6] ✓ 0.5s

```

1 model = KMeans(n_clusters=3)
2 plt.figure(figsize=(5,3),dpi=100)
3 display_categories(model,blobs)


```

4] ✓ 0.8s



```
1 model = KMeans(n_clusters=2)
2 plt.figure(figsize=(5,3),dpi=100)
3 display_categories(model,blobs)
```

5] ✓ 0.6s




```

1 model = KMeans(n_clusters=2)
2 plt.figure(figsize=(5,3),dpi=100)
3 display_categories(model,moons)

```

✓ 0.5s




```

1 model = KMeans(n_clusters=6)
2 plt.figure(figsize=(5,3),dpi=100)
3 display_categories(model,moons)


```

✓ 0.9s



```
1 model = KMeans(n_clusters=2)
2 plt.figure(figsize=(5,3),dpi=100)
3 display_categories(model,circles)
```

✓ 0.5s




DBSCAN

```
1 from sklearn.cluster import DBSCAN
```

✓ 0.1s

```
1 model = DBSCAN()  
2 plt.figure(figsize=(5,3),dpi=100)  
3 display_categories(model,blobs)
```

✓ 0.9s




```

1 model = DBSCAN(eps=0.15)
2 plt.figure(figsize=(5,3),dpi=100)
3 display_categories(model,moons)

```

✓ 0.4s




```

1 model = DBSCAN(eps=0.15)
2 plt.figure(figsize=(5,3),dpi=100)
3 display_categories(model,circles)

```

✓ 0.6s



▼ Hyperparameters

DBSCAN Hyperparameters

Hyperparameter Examples

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
```

✓ 0.2s

```
1 two_blobs = pd.read_csv('cluster_two_blobs.csv')
2 two_blobs_outliers = pd.read_csv('cluster_two_blobs_outliers.csv')
```

✓ 0.1s

```
1 two_blobs
```


✓ 0.1s

	X1	X2
0	0.046733	1.765120
1	-8.994134	-6.508186
2	0.650539	1.264533

```
1 sns.scatterplot(data=two_blobs, x="X1", y="X2")
```

```
4] ✓ 0.4s
```

```
<AxesSubplot:xlabel='X1', ylabel='X2'>
```



```
1 two_blobs_outliers.head()
```


✓ 0.7s

	X1	X2
0	0.046733	1.765120
1	-8.994134	-6.508186
2	0.650539	1.264533
3	-9.501554	-6.736493
4	0.057050	0.188215

```
1 sns.scatterplot(data=two_blobs_outliers, x="X1", y="X2")
```

✓ 0.4s

<AxesSubplot:xlabel='X1', ylabel='X2'>



```
1 def display_categories(model,data):
2
3     labels = model.fit_predict(data)
4     sns.scatterplot(data=data, x="X1", y="X2", hue=labels, palette="Set1")
```

✓ 0.6s

DBSCAN

```
1 from sklearn.cluster import DBSCAN
```


✓ 0.9s

```
1 dbscan = DBSCAN()
```

✓ 0.8s


```
1 display_categories(dbscan, two_blobs)
```

✓ 0.5s



1 `display_categories(dbscan, two_blobs_outliers)`

1] ✓ 0.4s



Epsilon

`eps : float, default=0.5`

| The maximum distance between two samples for one to be considered

| as in the neighborhood of the other. This is not a maximum bound

| on the distances of points within a cluster. This is the most


| important DBSCAN parameter to choose appropriately for your data set

| and distance function.

```
1 dbscan = DBSCAN(eps=0.001)
2 display_categories(dbscan, two_blobs_outliers)
```


✓ 0.4s

Python





```
1 dbscan = DBSCAN(eps=10)
2 display_categories(dbscan,two_blobs_outliers)
```

✓ 0.4s



```
1 dbscan = DBSCAN(eps=1)
2 display_categories(dbscan,two_blobs_outliers)
```

✓ 0.5s



```
1 np.sum(dbscan.labels_ == -1)
✓ 0.7s

3

1 100 * np.sum(dbscan.labels_ == -1) / len(dbscan.labels_)
✓ 0.1s


0.29910269192422734
```

```
1 outlier_percent = []
2 number_of_outliers = []
3
4 for eps in np.linspace(0.001,10,100):
5     # Create Model
6     dbscan = DBSCAN(eps=eps)
7     dbscan.fit(two_blobs_outliers)
8
9     # Log Number of Outliers
10    number_of_outliers.append(np.sum(dbscan.labels_ == -1))
11
12    # Log percentage of points that are outliers
13    perc_outliers = 100 * np.sum(dbscan.labels_ == -1) / len(dbscan.labels_)
14
15    outlier_percent.append(perc_outliers)
16
17
✓ 1.8s
```

```
1 sns.lineplot(x=np.linspace(0.001,10,100), y=number_of_outliers)
2 plt.xlim(0,1)
```

✓ 0.4s


(0.0, 1.0)



```
1 sns.lineplot(x=np.linspace(0.001,10,100), y=outlier_percent)
2 plt.xlim(0,1)
```

✓ 0.3s


(0.0, 1.0)



```
1 sns.lineplot(x=np.linspace(0.001,10,100), y=number_of_outliers)
2 plt.xlim(0,1)
3 plt.ylim(0,10)
4 plt.xlim(0,2)
5 plt.hlines(y=3,xmin=0,xmax=2,colors='red',ls='--')
```

✓ 0.3s

<matplotlib.collections.LineCollection at 0x251d8d62b50>




```

1 sns.lineplot(x=np.linspace(0.001,10,100),y=number_of_outliers)
2 plt.ylim(0,10)
3 plt.xlim(0,6)
4 plt.hlines(y=3,xmin=0,xmax=10,colors='red',ls='--')

```

1] ✓ 0.3s

<matplotlib.collections.LineCollection at 0x251d8dfb070>



Minimum Samples

```

| min_samples : int, default=5
|   The number of samples (or total weight) in a neighborhood for a point
|   to be considered as a core point. This includes the point itself.

```

How to choose minimum number of points?

<https://stats.stackexchange.com/questions/88872/a-routine-to-choose-eps-and-minpts-for-dbscan>

```

1 outlier_percent = []
2
3 for n in np.arange(1,100):
4
5     # Create Model
6     dbscan = DBSCAN(min_samples=n)
7     dbscan.fit(two_blobs_outliers)
8
9     # Log percentage of points that are outliers
10    perc_outliers = 100 * np.sum(dbscan.labels_ == -1) / len(dbscan.labels_)
11
12    outlier_percent.append(perc_outliers)
13


```

2] ✓ 1.8s

```
1 sns.lineplot(x=np.arange(1,100),y=outlier_percent)
```

✓ 0.3s

<AxesSubplot:>




```

1 num_dim = two_blobs_outliers.shape[1]
2
3 dbscan = DBSCAN(min_samples=2*num_dim)
4 display_categories(dbscan,two_blobs_outliers)

```

✓ 0.5s




```

1 num_dim = two_blobs_outliers.shape[1]
2
3 dbscan = DBSCAN(eps=0.75,min_samples=2*num_dim)
4 display_categories(dbscan,two_blobs_outliers)


```

✓ 0.5s




```
1 dbscan = DBSCAN(min_samples=1)
2 display_categories(dbscan,two_blobs_outliers)
```

✓ 0.9s



```
1 dbscan = DBSCAN(eps=0.75,min_samples=1)
2 display_categories(dbscan,two_blobs_outliers)
```

✓ 0.6s



▼ Customer Data Project

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
```

✓ 2.7s

```
1 df = pd.read_csv('wholesome_customers_data.csv')
```

✓ 0.7s

```
1 df
```

✓ 0.4s

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen
0	2	3	12669	9656	7561	214	2674	1338
1	2	3	7057	9810	9568	1762	3293	1776
2	2	3	6353	8808	7684	2405	3516	7844
3	1	3	13265	1196	4221	6404	507	1788
4	2	3	22615	5410	7198	3915	1777	5185
...

```
1 df.info()
```

✓ 0.1s

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 440 entries, 0 to 439
```

```
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
0	Channel	440 non-null	int64
1	Region	440 non-null	int64
2	Fresh	440 non-null	int64
3	Milk	440 non-null	int64
4	Grocery	440 non-null	int64
5	Frozen	440 non-null	int64
6	Detergents_Paper	440 non-null	int64
7	Delicassen	440 non-null	int64

```
dtypes: int64(8)
```


```
memory usage: 27.6 KB
```

EDA

```
1 sns.scatterplot(data=df, x="Milk", y="Grocery", hue="Channel")
```

✓ 0.5s


<AxesSubplot:xlabel='Milk', ylabel='Grocery'>



```
1 sns.histplot(data=df, x="Milk", hue="Channel", palette="Set1", multiple="stack")
```

✓ 0.5s

<AxesSubplot:xlabel='Milk', ylabel='Count'>




```
1 clustermap(df.drop(["Region","Channel"], axis=1).corr(), annot=True, row_cluster=False)
```

✓ 1.5s

Python

<seaborn.matrix.ClusterGrid at 0x1c5b326d3d0>

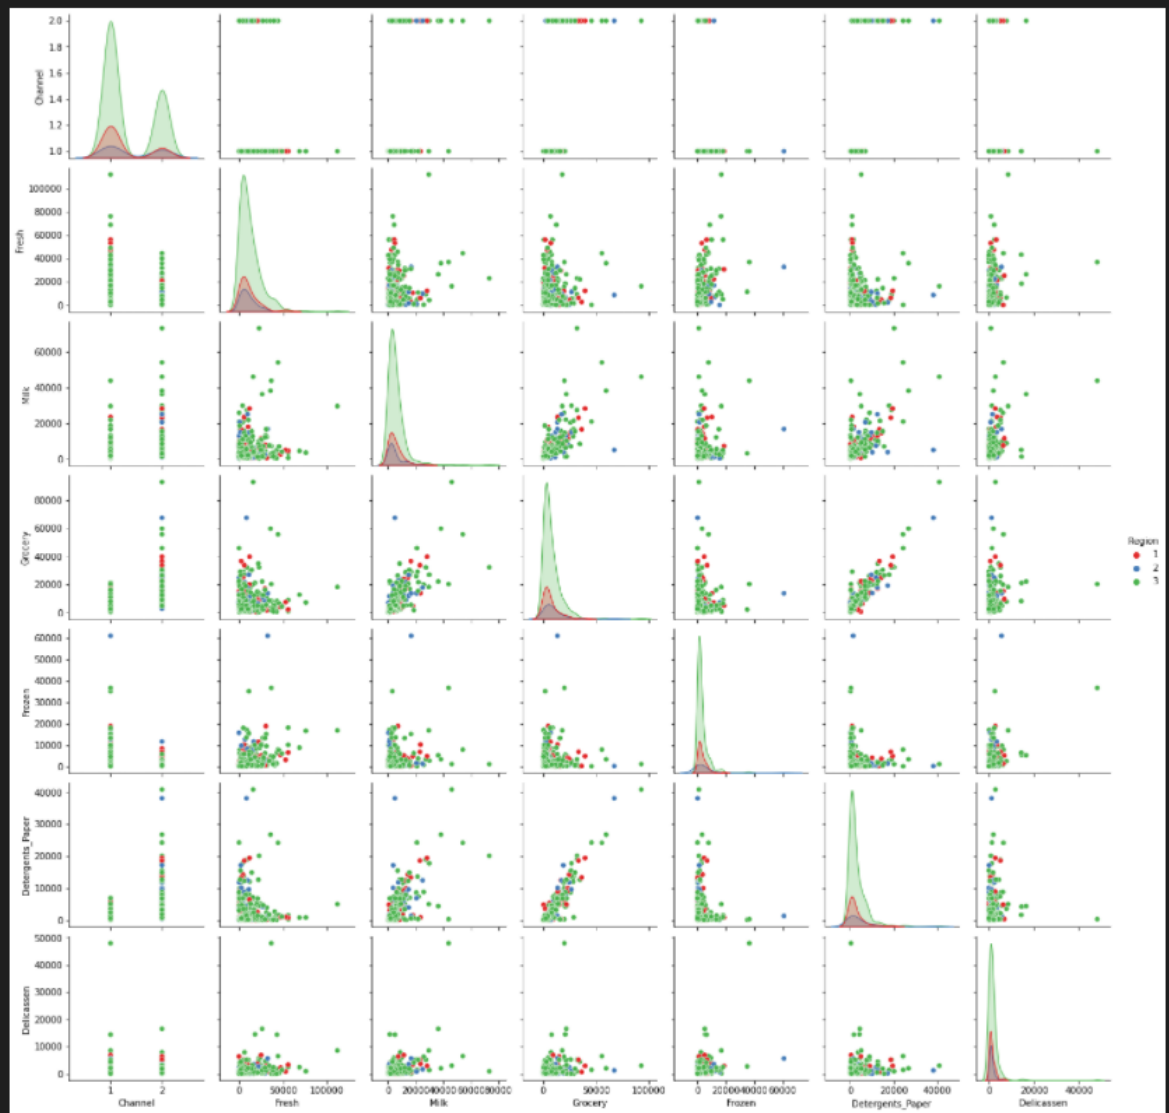


```
1 sns.pairplot(df, hue="Region", palette="Set1")
```

✓ 18.7s

Python

<seaborn.axisgrid.PairGrid at 0x1c5b37bbac0>



ML Model

```
1 from sklearn.preprocessing import StandardScaler
2 scaler = StandardScaler()
```

✓ 0.2s

```
1 scaled_X = scaler.fit_transform(df)
```

✓ 0.1s

```
1 from sklearn.cluster import DBSCAN
```

✓ 0.7s


```
1 outlier_percent = []
2
3 for eps in np.linspace(0.001,3,50):
4
5     dbscan = DBSCAN(eps=eps, min_samples=2*scaled_X.shape[1])
6     dbscan.fit(scaled_X)
7
8     perc_outliers = 100*np.sum(dbscan.labels_== -1) / len(dbscan.labels_)
9
10    outlier_percent.append(perc_outliers)
```

✓ 0.8s

```
1 sns.lineplot(x=np.linspace(0.001,3,50), y=outlier_percent)
```

✓ 0.3s

<AxesSubplot:>



```
1 dbscan = DBSCAN(eps=2, min_samples=scaled_X.shape[1])
```

✓ 0.1s

```
1 dbscan.fit(scaled_X)
```


✓ 0.8s

```
DBSCAN(eps=2, min_samples=8)
```

```
1 sns.scatterplot(data=df,x='Grocery',y='Milk',hue=dbscan.labels_)
```

✓ 0.7s


<AxesSubplot:xlabel='Grocery', ylabel='Milk'>



```
1 sns.scatterplot(data=df,x='Detergents_Paper',y='Milk',hue=dbscan.labels_)
```

✓ 0.8s

<AxesSubplot:xlabel='Detergents_Paper', ylabel='Milk'>




```
1 df["Labels"] = dbscan.labels_
✓ 0.1s
```

```
1 df
✓ 0.1s
```

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen	Labels
0	2	3	12669	9656	7561	214	2674	1338	0
1	2	3	7057	9810	9568	1762	3293	1776	0
2	2	3	6353	8808	7684	2405	3516	7844	0
3	1	3	13265	1196	4221	6404	507	1788	1
4	2	3	22615	5410	7198	3915	1777	5185	0
...
435	1	3	29703	12051	16027	13135	182	2204	1
436	1	3	39228	1431	764	4510	93	2346	1
437	2	3	14531	15488	30243	437	14841	1867	0
438	1	3	10290	1981	2232	1038	168	2125	1
439	1	3	2787	1698	2510	65	477	52	1

440 rows x 9 columns

```
1 cats = df.drop(["Channel","Region"], axis=1)
2 cat_mean = cats.groupby("Labels").mean()
3 cat_mean
✓ 0.6s
```

	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen
Labels						
-1	28678.285714	24176.523810	28797.857143	11535.000000	11932.523810	7367.380952
0	8134.862595	8909.916031	14004.427481	1450.595420	6080.832061	1533.519084
1	12542.430556	3039.760417	3677.871528	3192.315972	766.267361	1094.920139

```


1 plt.figure(figsize=(8,4), dpi=120)
2 sns.heatmap(cat_mean, annot=True)

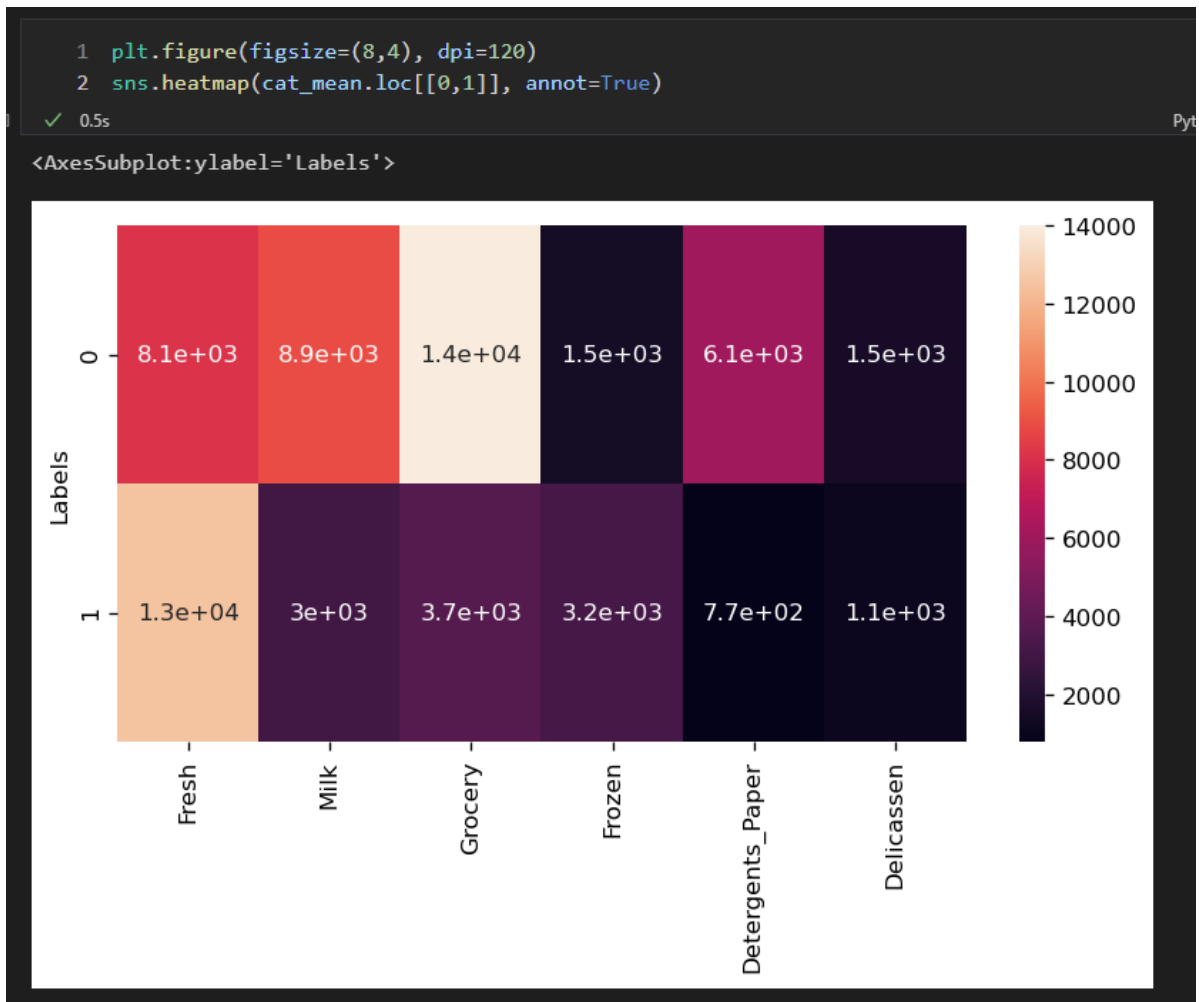
```

✓ 0.7s

Py

<AxesSubplot:ylabel='Labels'>





```


1 from sklearn.preprocessing import MinMaxScaler
✓ 0.5s

1 scaler = MinMaxScaler()
2 data=scaler.fit_transform(cat_mean)
3 scaled_cat = pd.DataFrame(data, cat_mean.index, cat_mean.columns)
✓ 0.2s

1 sns.heatmap(pd.DataFrame(data, cat_mean.index, cat_mean.columns))
✓ 0.4s

```

<AxesSubplot:ylabel='Labels'>



```
1 sns.heatmap(scaled_cat.loc[[0,1]],annot=True)
```

✓ 0.5s

<AxesSubplot:ylabel='Labels'>

