

# Lecture 3 Classification Techniques

May 11, 2022

# 1 Classification Workbook

Welcome to the Classification workbook. Let's start by importing some of our custom libraries.

```
[1]: import numpy as np
import pandas as pd

from matplotlib import pyplot as plt
%matplotlib inline

import seaborn as sns
sns.set_style('darkgrid')
```

## 1.1 1. Classification Algorithms

In the class, we discussed five different classification methods. Let's investigate them in more detail here.

### 1.1.1 1.1. Binary Classification

Binary classification is when the target is determined with True or False. Let's create a simple x and y arrays for exploring conditional classification.

```
[2]: x = np.linspace(0, 1, 100)
```

```
[3]: np.random.seed(123)
noise = np.random.uniform(-0.2, 0.2, 100)
```

```
[4]: y = ((x+noise)>0.5).astype(int)
```


[5] : y

```
[5]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,  
          1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0,  
          1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
          1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1])
```

```
[6]: fig = plt.figure(figsize=(6,4))
ax = plt.subplot(111)

ax.plot(x, x>0.5, 'k.')
ax.scatter(x, y, color='b')
```

```
[6]: <matplotlib.collections.PathCollection at 0x13d9037c0>
```



Simple as that, we built a mask that attributes values below 0.5 to 0 and above 0.5 to 1. When we use an encoder for a binary classification problem, we can use this type of classification easily.

### 1.1.2 1.2. Logistic Regression

We have used Logistic Regression before, but unlike the name implies LogisticRegression is best suited for classification problems.

```
[7]: from sklearn.linear_model import LogisticRegression
```

```
[8]: lr_model = LogisticRegression()
```

```
[9]: lr_model.fit(x.reshape(100,1),y)
```

```
[9]: LogisticRegression()
```

```
[10]: lr_model.predict(x.reshape(100,1))
```

```
[10]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
            0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
            0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
            1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
            1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1])
```

```
[11]: (lr_model.predict(x.reshape(100,1))==y)[:40]
```

```
[11]: array([ True,  True,  True,  True,  True,  True,  True,  True,  True,
            True,  True,  True,  True,  True,  True,  True,  True,  True,
            True,  True,  True,  True,  True,  True,  True,  True,  True,
            True,  True,  True,  True,  True,  True,  True,  True,  True,
            True, False, False,  True])
```

This tells us model doesn't always predict the values of y. Remember with LogisticRegression, it is always wise to look at the probability.

```
[12]: lr_model.predict_proba(x.reshape(100,1))[:40]
```

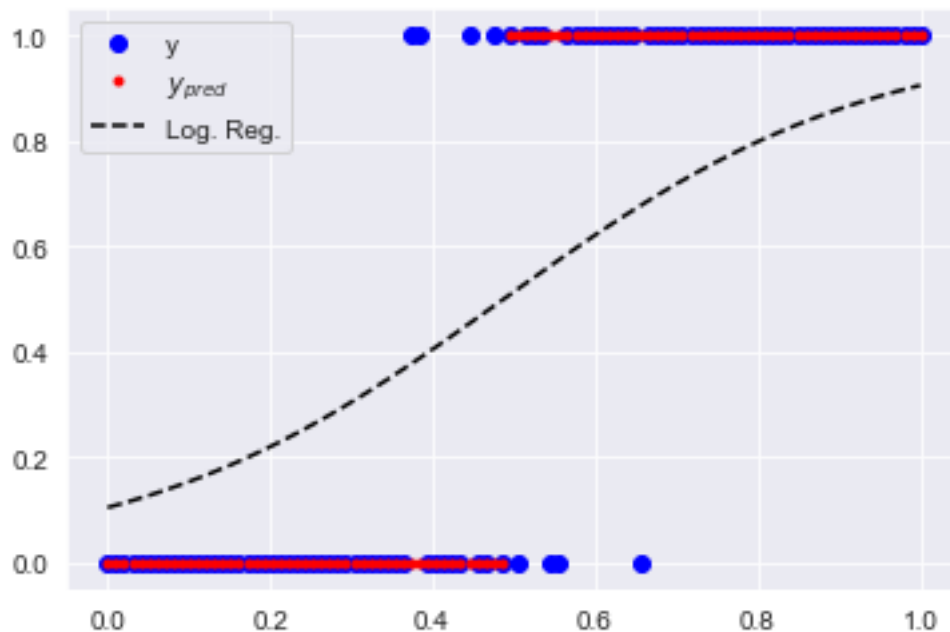
```
[12]: array([[0.89531454, 0.10468546],
            [0.89107222, 0.10892778],
            [0.88667974, 0.11332026],
            [0.88213356, 0.11786644],
            [0.87743021, 0.12256979],
            [0.87256629, 0.12743371],
            [0.8675385 , 0.1324615 ],
            [0.86234363, 0.13765637],
            [0.85697862, 0.14302138],
            [0.85144054, 0.14855946],
            [0.84572661, 0.15427339],
            [0.83983425, 0.16016575],
            [0.83376108, 0.16623892],
            [0.82750492, 0.17249508],
            [0.82106385, 0.17893615],
            [0.8144362 , 0.1855638 ],
            [0.80762059, 0.19237941],
            [0.80061592, 0.19938408],
            [0.79342145, 0.20657855],
            [0.78603676, 0.21396324],
            [0.77846179, 0.22153821],
            [0.77069688, 0.22930312],
            [0.76274275, 0.23725725],
            [0.75460057, 0.24539943],
            [0.74627191, 0.25372809],
            [0.73775882, 0.26224118],
            [0.72906379, 0.27093621],
            [0.72018981, 0.27981019],
```

```
[0.71114034, 0.28885966],
[0.70191933, 0.29808067],
[0.69253123, 0.30746877],
[0.682981  , 0.317019  ],
[0.67327408, 0.32672592],
[0.66341641, 0.33658359],
[0.65341443, 0.34658557],
[0.64327504, 0.35672496],
[0.63300563, 0.36699437],
[0.62261402, 0.37738598],
[0.61210848, 0.38789152],
[0.60149767, 0.39850233]])
```

```
[13]: fig = plt.figure(figsize=(6,4))
      ax = plt.subplot(111)

      ax.plot(x, y, 'bo', label='y')
      ax.plot(x, lr_model.predict(x.reshape(100,1)), 'r.', label='$y_{pred}$')
      ax.plot(x, lr_model.predict_proba(x.reshape(100,1))[:,1], 'k--', label='Log. Prob.
      ↪Reg.')
      ax.legend()
```

```
[13]: <matplotlib.legend.Legend at 0x13e26ba00>
```



The probabilities are not as “certain”. Remember, we can always fine tune the logistic regression with C, penalty parameter.

### 1.1.3 1.3. Decision Tree and Random Forest Classification

This is also an old friend of ours. Let's see how we can use RandomForest as a classifier.

```
[14]: from sklearn.ensemble import RandomForestClassifier
```

```
[15]: rf_model = RandomForestClassifier(random_state=123)
```

```
[16]: rf_model.fit(x.reshape(100,1),y)
```

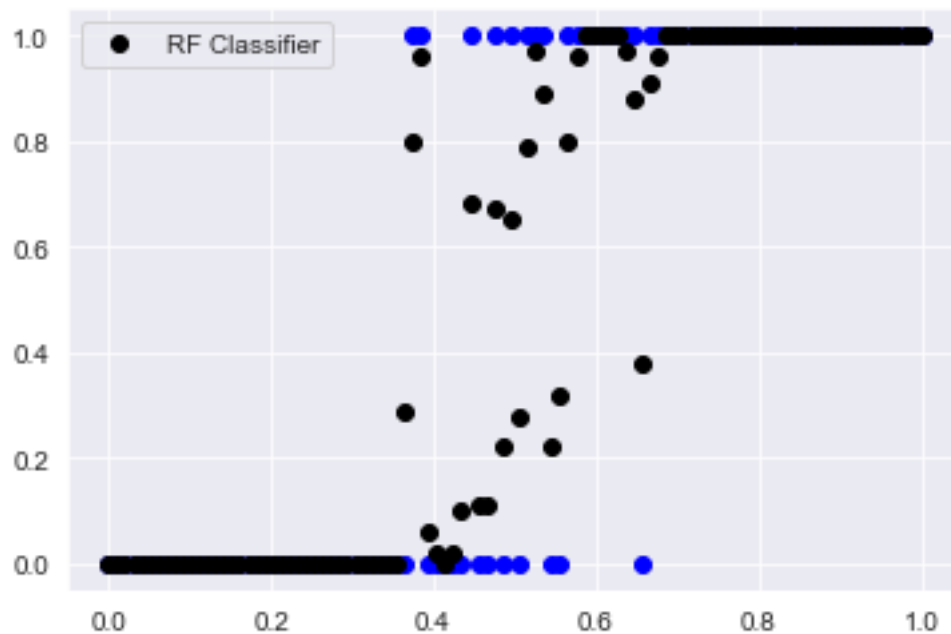
```
[16]: RandomForestClassifier(random_state=123)
```

```
[17]: X = x.reshape(100,1)
```

```
[18]: pred = rf_model.predict_proba(X)
```

```
[19]: fig = plt.figure(figsize=(6,4))  
ax = plt.subplot(111)  
  
ax.scatter(x, y, color='b')  
ax.plot(X, pred[:,1], 'ko', label='RF Classifier')  
ax.legend()
```

```
[19]: <matplotlib.legend.Legend at 0x13e7abd60>
```



Not a great fit, but luckily we have learnt how to tune Random Forests.

```
[20]: rf_model2 = RandomForestClassifier(random_state=123, n_estimators=10,
    ↪min_samples_leaf=20)

[21]: rf_model2.fit(X,y)

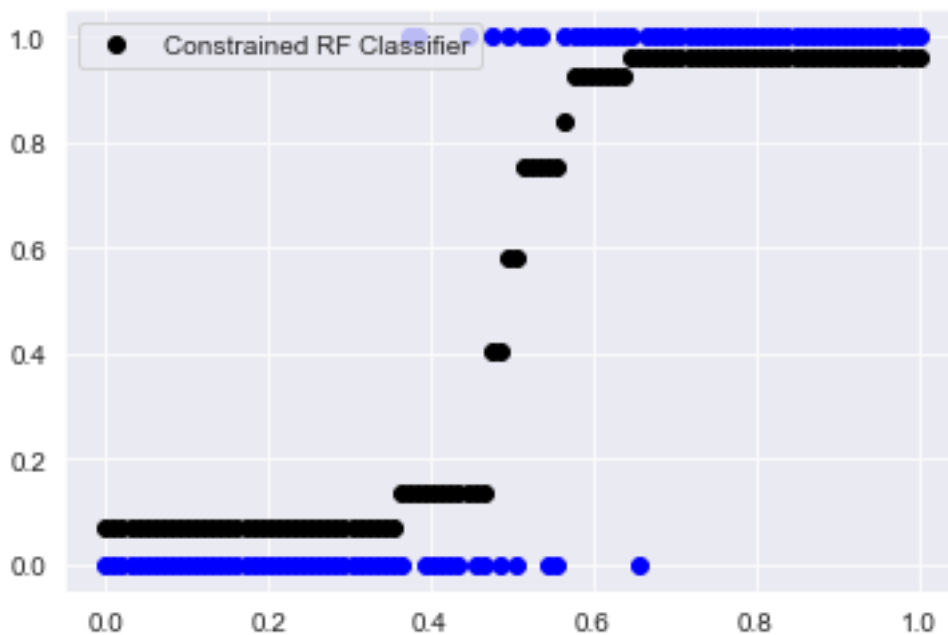
[21]: RandomForestClassifier(min_samples_leaf=20, n_estimators=10, random_state=123)

[22]: pred2 = rf_model2.predict_proba(X)

[23]: fig = plt.figure(figsize=(6,4))
    ax = plt.subplot(111)

    ax.scatter(x, y, color='b')
    ax.plot(X, pred2[:,1], 'ko', label='Constrained RF Classifier')
    ax.legend(loc=2)

[23]: <matplotlib.legend.Legend at 0x13e825f70>
```



#### 1.1.4 1.4. k-Nearest Neighbour Classification

```
[24]: from sklearn.neighbors import KNeighborsClassifier

[25]: knn_model = KNeighborsClassifier(n_neighbors=3)


[26]: knn_model.fit(X,y)
```

```
[26]: KNeighborsClassifier(n_neighbors=3)
```

```
[27]: pred4 = knn_model.predict_proba(X)
```

```
[28]: fig = plt.figure(figsize=(6,4))  
ax = plt.subplot(111)  
  
ax.scatter(x, y, color='b')  
ax.plot(X, pred4[:,1], 'r.', label='kNN')  
ax.legend(loc=2)
```

```
[28]: <matplotlib.legend.Legend at 0x13d951df0>
```



Again, we can finetune the kNN to get better results.

```
[29]: knn_model2 = KNeighborsClassifier(n_neighbors=10, weights='uniform')  
knn_model2.fit(X,y)
```


```
[29]: KNeighborsClassifier(n_neighbors=10)
```

```
[30]: pred5 = knn_model2.predict_proba(X)
```

```
[31]: fig = plt.figure(figsize=(6,4))  
ax = plt.subplot(111)  
  
ax.scatter(x, y, color='b')
```

```
ax.plot(X, pred5[:,1], 'r.', label='kNN')
ax.legend(loc=2)
```

[31]: <matplotlib.legend.Legend at 0x13e8cc790>



## 1.2 2. Evaluation Methods

Let's remember what our x and y were.

[32]: `print(X, y)`

```
[[0.      ]
 [0.01010101]
 [0.02020202]
 [0.03030303]
 [0.04040404]
 [0.05050505]
 [0.06060606]
 [0.07070707]
 [0.08080808]
 [0.09090909]
 [0.1010101 ]
 [0.11111111]
 [0.12121212]
 [0.13131313]
 [0.14141414]
```



[0.15151515]  
[0.16161616]  
[0.17171717]  
[0.18181818]  
[0.19191919]  
[0.2020202 ]  
[0.21212121]  
[0.22222222]  
[0.23232323]  
[0.24242424]  
[0.25252525]  
[0.26262626]  
[0.27272727]  
[0.28282828]  
[0.29292929]  
[0.3030303 ]  
[0.31313131]  
[0.32323232]  
[0.33333333]  
[0.34343434]  
[0.35353535]  
[0.36363636]  
[0.37373737]  
[0.38383838]  
[0.39393939]  
[0.4040404 ]  
[0.41414141]  
[0.42424242]  
[0.43434343]  
[0.44444444]  
[0.45454545]  
[0.46464646]  
[0.47474747]  
[0.48484848]  
[0.49494949]  
[0.50505051]  
[0.51515152]  
[0.52525253]  
[0.53535354]  
[0.54545455]  
[0.55555556]  
[0.56565657]  
[0.57575758]  
[0.58585859]  
[0.5959596 ]  
[0.60606061]  
[0.61616162]  
[0.62626263]



```
[35]: pred2 = rf_model.predict(X)
```

```
[36]: pred3 = rf_model2.predict(X)
```

```
[37]: pred4 = knn_model.predict(X)
```

```
[38]: print( 'Logistic Reg. Acc:', accuracy_score(y, pred))
print( 'Random Forest Acc:', accuracy_score(y, pred2))
print( 'Constrained Random Forest Acc:', accuracy_score(y, pred3))
print( 'KNN Acc:', accuracy_score(y, pred4))
```

Logistic Reg. Acc: 0.92

Random Forest Acc: 1.0

Constrained Random Forest Acc: 0.92

KNN Acc: 0.94

This is what the metric says, but is something fishy?

### 1.2.2 2.2. Confusion Matrix

```
[39]: from sklearn.metrics import confusion_matrix
```

```
[40]: print('Log. Reg.:\n', confusion_matrix(y, pred))
```

Log. Reg.:

[[45 4]

[ 4 47]]

```
[41]: print('Random Forest:\n', confusion_matrix(y, pred2))
```

Random Forest:

[[49 0]

[ 0 51]]

```
[42]: print('Const. RF.:\n', confusion_matrix(y, pred3))
```

Const. RF.:

[[45 4]

[ 4 47]]

```
[43]: print('KNN:\n', confusion_matrix(y, pred4))
```

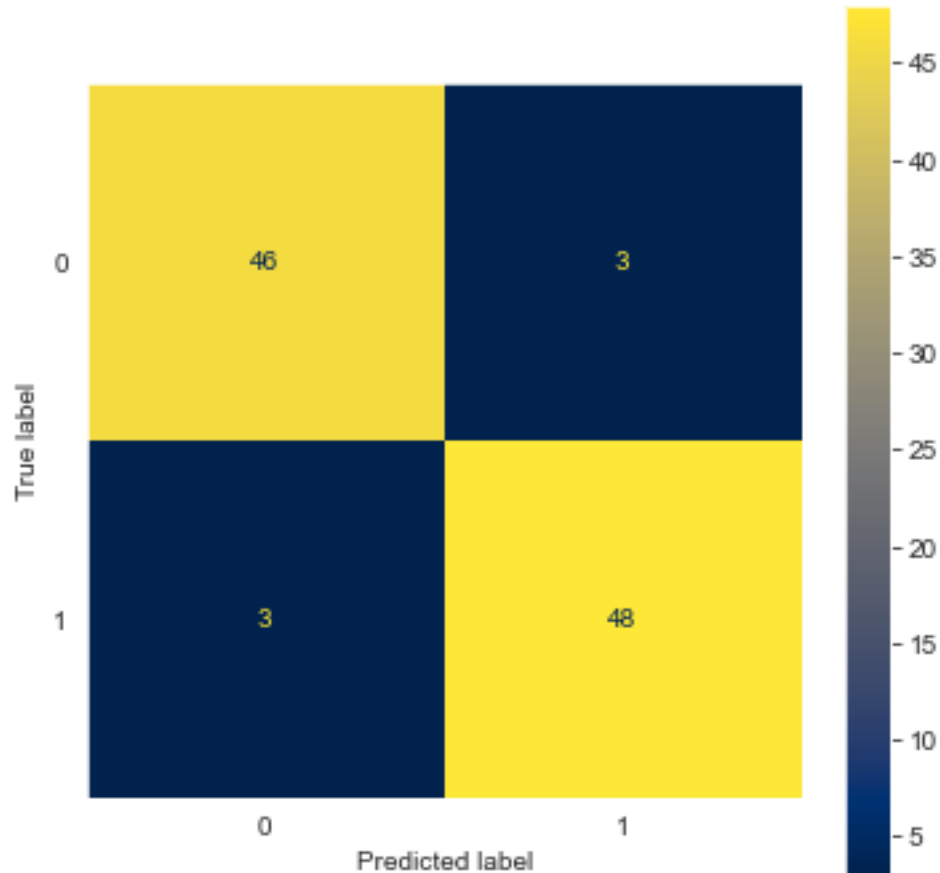
KNN:

[[46 3]

[ 3 48]]

```
[44]: from sklearn.metrics import plot_confusion_matrix
fig = plt.figure(figsize=[6,6])
```

```
ax = plt.subplot(111)
cb = plot_confusion_matrix(knn_model, X, y, ax=ax, cmap='cividis')
ax.grid()
plt.show()
```



Okay, this is more revealing about the performance of the individual model.

### 1.2.3 2.3. Precision-Recall

Precision and recall is another way of assessing the performance of a model.

```
[45]: from sklearn.metrics import precision_recall_curve
```

```
[46]: print('Log. Reg.:\n', precision_recall_curve(y, pred))
      print('Random Forest:\n', precision_recall_curve(y, pred2))
      print('Const. RF.:\n', precision_recall_curve(y, pred3))
      print('KNN:\n', precision_recall_curve(y, pred4))
```

```
Log. Reg.:
(array([0.51          , 0.92156863, 1.          ]), array([1.          , 0.92156863, 0.          ]))
```

```
]), array([0, 1]))
```

Random Forest:

```
(array([1., 1.]), array([1., 0.]), array([1]))
```

Const. RF.:

```
(array([0.51      , 0.92156863, 1.      ]), array([1.      , 0.92156863, 0.      ]), array([0, 1]))
```

KNN:


```
(array([0.51      , 0.94117647, 1.      ]), array([1.      , 0.94117647, 0.      ]), array([0, 1]))
```

Well, this didn't explain much.

```
[47]: from sklearn.metrics import PrecisionRecallDisplay
```

```
[48]: precision, recall, thresholds = precision_recall_curve(y, pred)
disp = PrecisionRecallDisplay(precision=precision, recall=recall)
disp.plot()
disp.ax_.set_title('Log. Reg.')
```


```
[48]: Text(0.5, 1.0, 'Log. Reg.')
```



```
[49]: precision, recall, thresholds = precision_recall_curve(y, pred)
plt.title("Precision-Recall vs Threshold Chart")
plt.plot(thresholds, precision[:-1], "bo", label="Precision")
plt.plot(thresholds, recall[:-1], "ro", label="Recall")
```


```
plt.ylabel("Precision, Recall")
plt.xlabel("Threshold")
plt.legend(loc="lower left")
plt.ylim([0,1.2])
```

[49]: (0.0, 1.2)



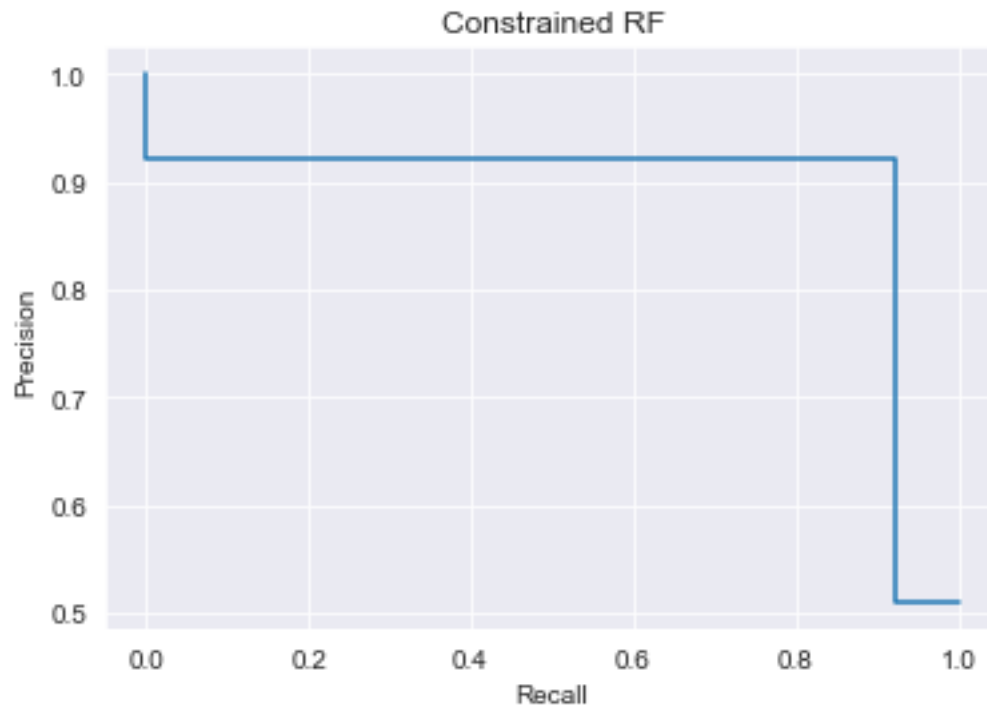
```
[50]: precision, recall, _ = precision_recall_curve(y, pred2)
disp = PrecisionRecallDisplay(precision=precision, recall=recall)
disp.plot()
disp.ax_.set_title('RF')
```

[50]: Text(0.5, 1.0, 'RF')



```
[51]: precision, recall, _ = precision_recall_curve(y, pred3)
      disp = PrecisionRecallDisplay(precision=precision, recall=recall)
      disp.plot()
      disp.ax_.set_title('Constrained RF')
```


```
[51]: Text(0.5, 1.0, 'Constrained RF')
```



```
[52]: precision, recall, _ = precision_recall_curve(y, pred4)
      disp = PrecisionRecallDisplay(precision=precision, recall=recall)
      disp.plot()
      disp.ax_.set_title('KNN')
```

```
[52]: Text(0.5, 1.0, 'KNN')
```





Now we know! Unconstrained Random Forest is overfitting.

#### 1.2.4 2.4. ROC Curve

The last performance metric for classifications we will learn in class today is the ROC curves. This is also very popular among data scientists.

```
[53]: from sklearn.metrics import roc_curve, roc_auc_score
```

```
[54]: # Calculate ROC curve from y and predictions
fpr, tpr, thresholds = roc_curve(y, pred)
lr_df = pd.DataFrame({'FPR': fpr, 'TPR' : tpr, 'Thresholds' : thresholds})
```

```
[55]: fpr, tpr, thresholds = roc_curve(y, pred2)
rf_df = pd.DataFrame({'FPR': fpr, 'TPR' : tpr, 'Thresholds' : thresholds})
```

```
[56]: fpr, tpr, thresholds = roc_curve(y, pred3)
crf_df = pd.DataFrame({'FPR': fpr, 'TPR' : tpr, 'Thresholds' : thresholds})
```

```
[57]: fpr, tpr, thresholds = roc_curve(y, pred4)
knn_df = pd.DataFrame({'FPR': fpr, 'TPR' : tpr, 'Thresholds' : thresholds})
```

```
[58]: lr_df.head()
```

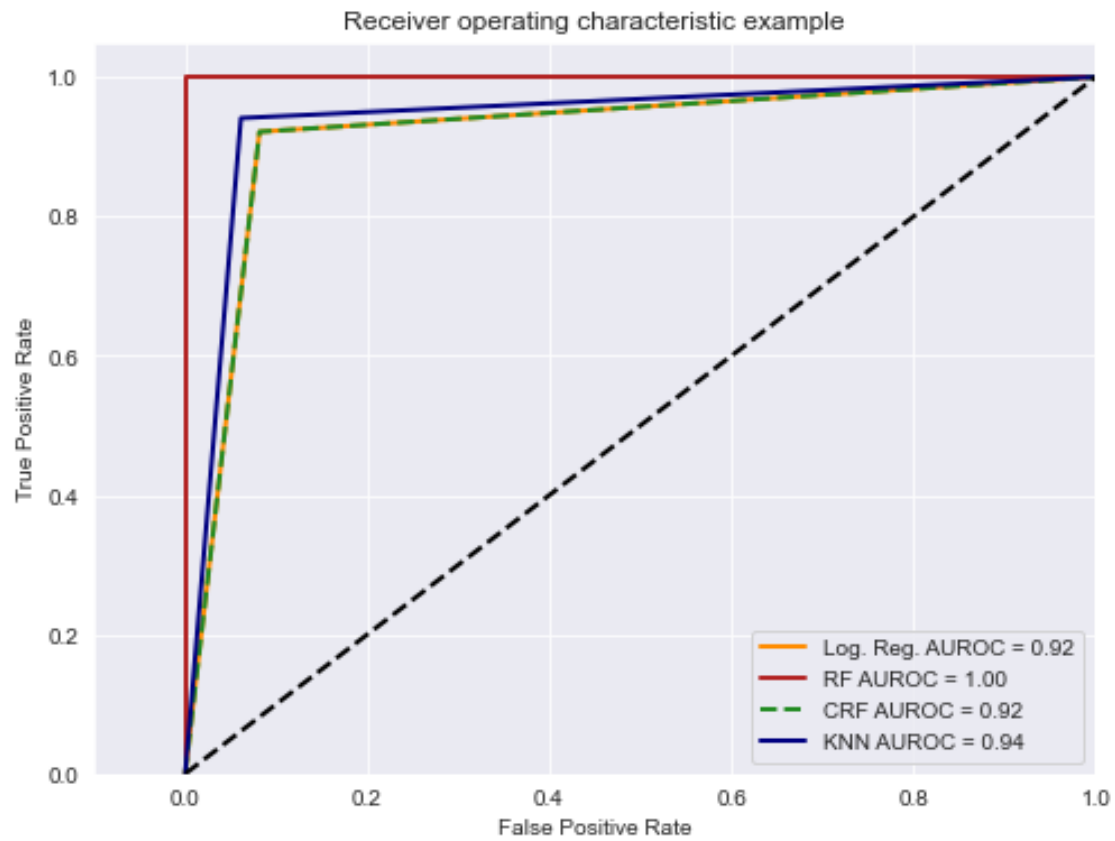
```
[58]:
```

	FPR	TPR	Thresholds
0	0.000000	0.000000	2
1	0.081633	0.921569	1
2	1.000000	1.000000	0

```
[59]: fig = plt.figure(figsize=(8,6))

ax = plt.subplot(111)
ax.plot(lr_df.FPR, lr_df.TPR, color='darkorange',
        linewidth=2, label='Log. Reg. AUROC = {:.2f}'.format(roc_auc_score(y,
↪pred)))
ax.plot(rf_df.FPR, rf_df.TPR, color='firebrick',
        linewidth=2, label='RF AUROC = {:.2f}'.format(roc_auc_score(y,
↪pred2)))
ax.plot(crf_df.FPR, crf_df.TPR, color='forestgreen', linestyle='--',
        linewidth=2, label='CRF AUROC = {:.2f}'.format(roc_auc_score(y,
↪pred3)))
ax.plot(knn_df.FPR, knn_df.TPR, color='navy',
        linewidth=2, label='KNN AUROC = {:.2f}'.format(roc_auc_score(y,
↪pred4)))

ax.plot([0, 1], [0, 1], color='k', lw=2, linestyle='--')
ax.set_xlim([-0.1, 1.0])
ax.set_ylim([0.0, 1.05])
ax.set_xlabel('False Positive Rate')
ax.set_ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
ax.legend(loc="lower right")
plt.show()
```



Congratulations, you have completed the Classification Workbook!

[ ]: